

Here Be Dragons

Writing Reliable Python Extensions in C

Paul Ross

Why?

- Blinding performance
- Leaner resources
- Interface with C/C++ libraries

Here Be Dragons



Here Be Dragons



Here Be Dragons



C

Here Be Dragons

CPython code!


C



Here Be Dragons

CPython code!

C



Automatic Memory Management

- Every Python object has a reference count
 - On creation this is set to 1
 - When it becomes 0 the object can be free'd

A PyObject

```
typedef struct _object {  
    Py_ssize_t ob_refcnt;  
    struct _typeobject *ob_type;  
} PyObject;
```

Reference Counting


```
>>> a = 'foo'  
>>> b = a  
>>> a = 'bar'  
>>> del a
```

Reference Counting

```
>>> a = 'foo'
```

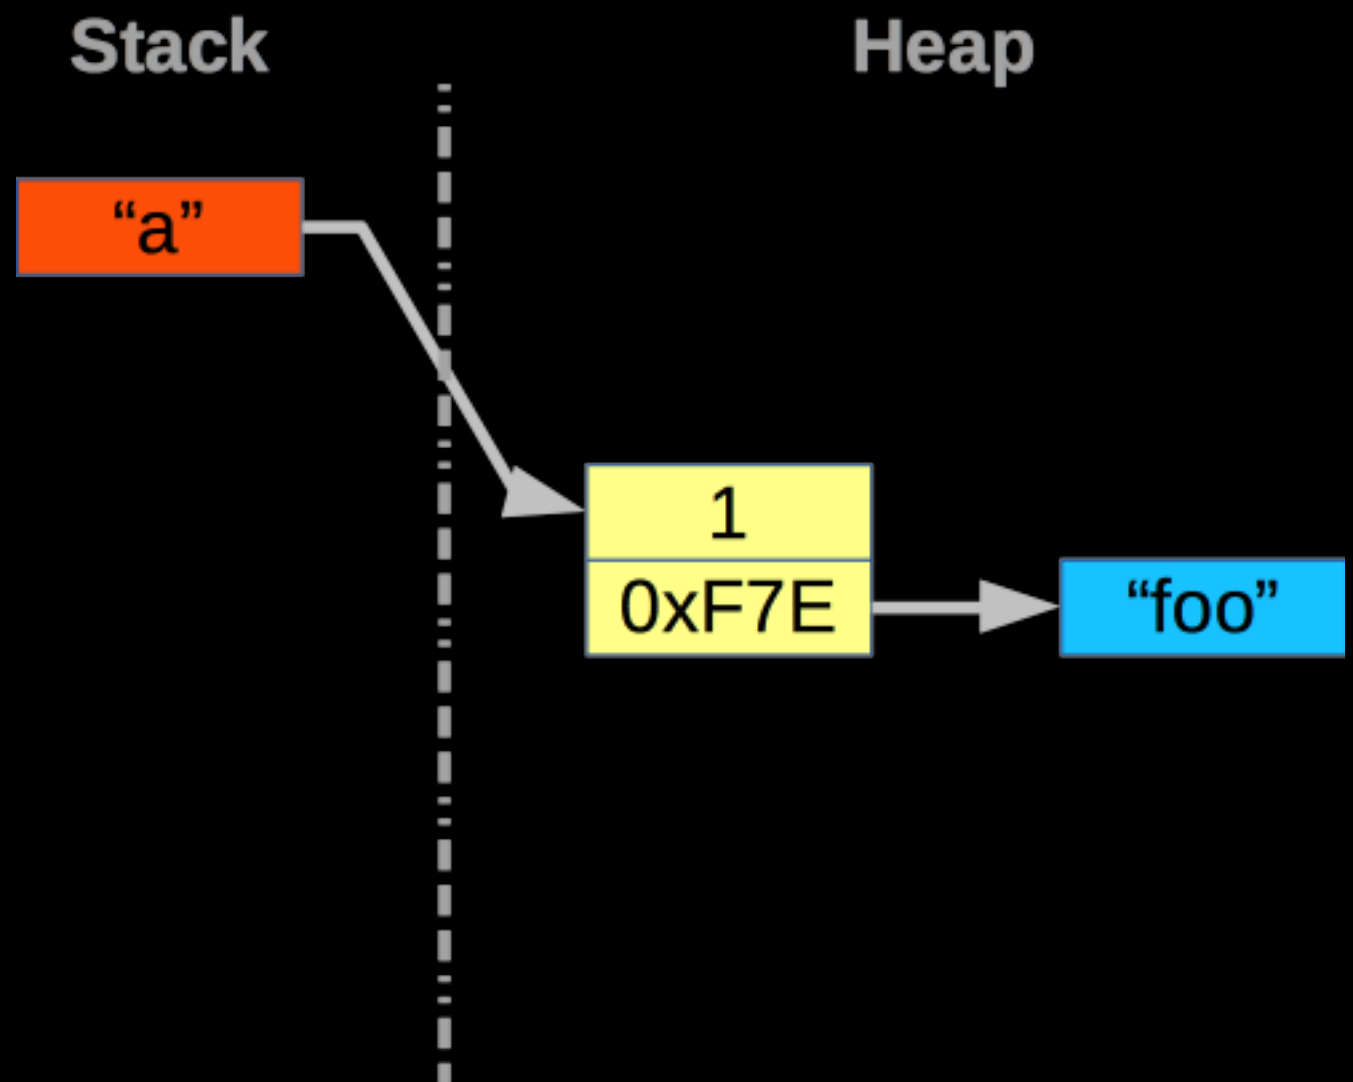
Reference Counting

```
>>> a = 'foo'
```




Reference Counting

```
>>> a = 'foo'  
>>> b = a
```



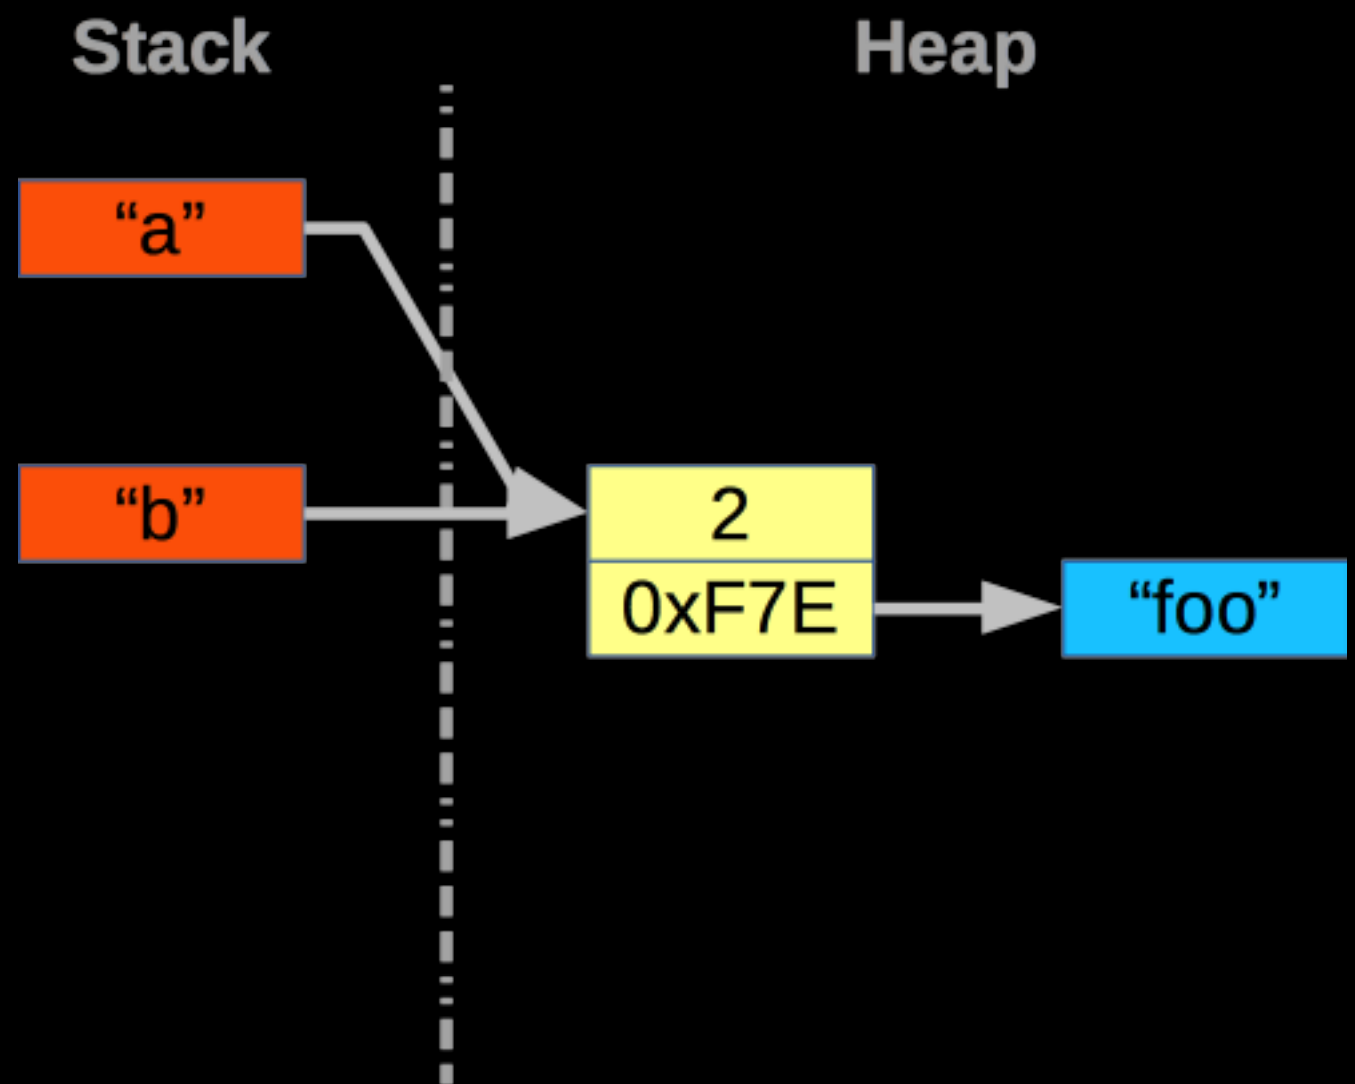
Reference Counting

```
>>> a = 'foo'  
>>> b = a
```



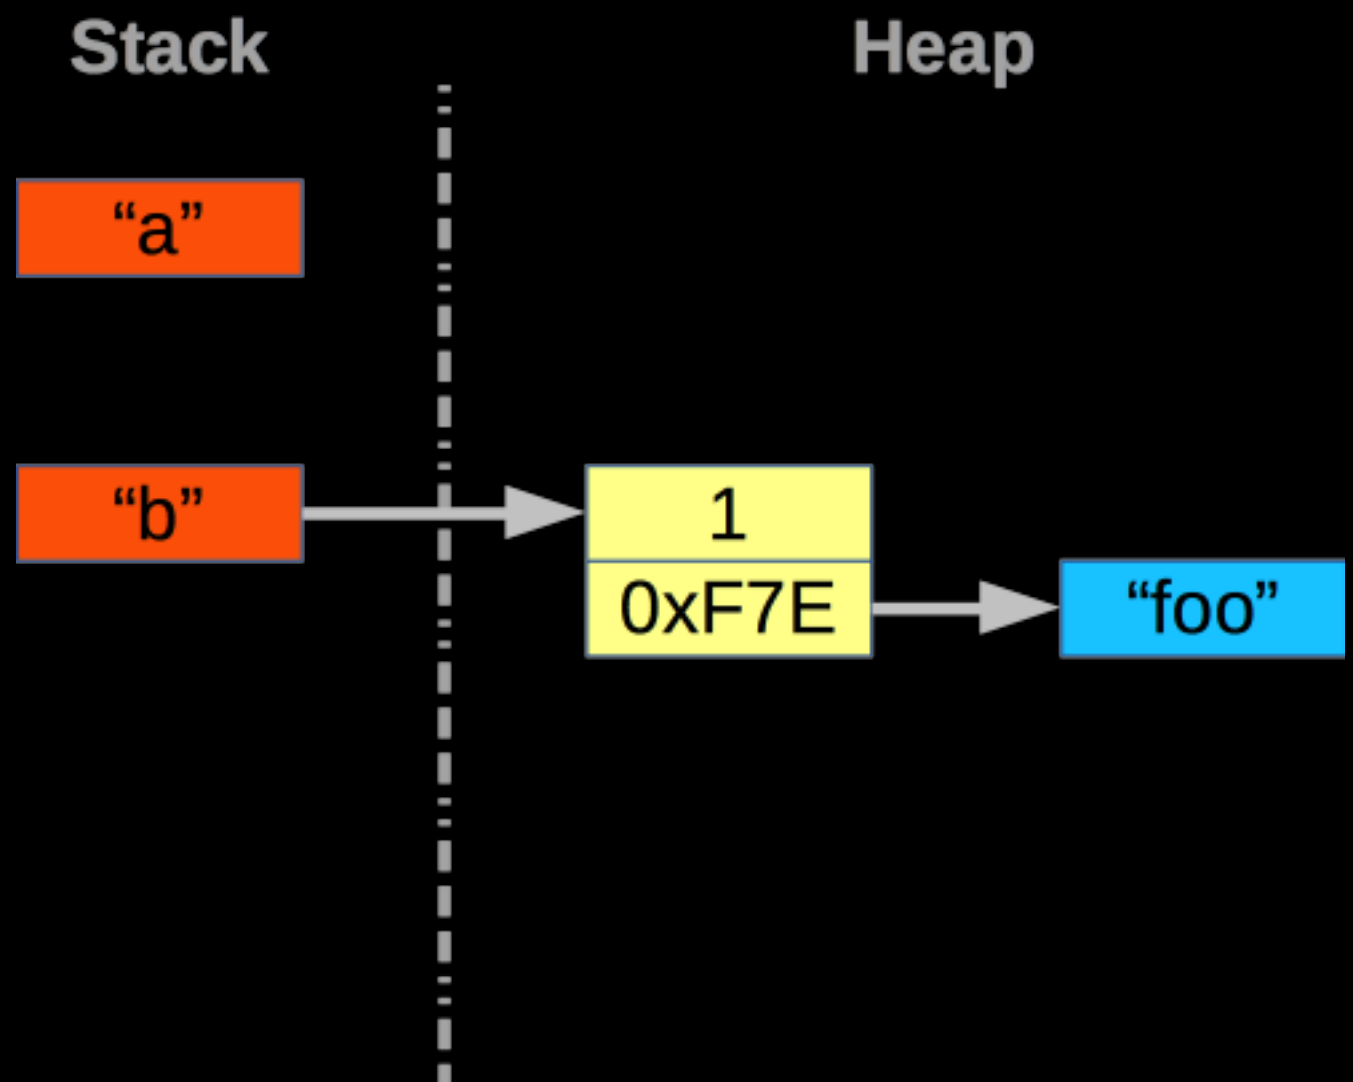
Reference Counting

```
>>> a = 'foo'  
>>> b = a  
>>> a = 'bar'
```




Reference Counting

```
>>> a = 'foo'  
>>> b = a  
>>> a = 'bar'
```



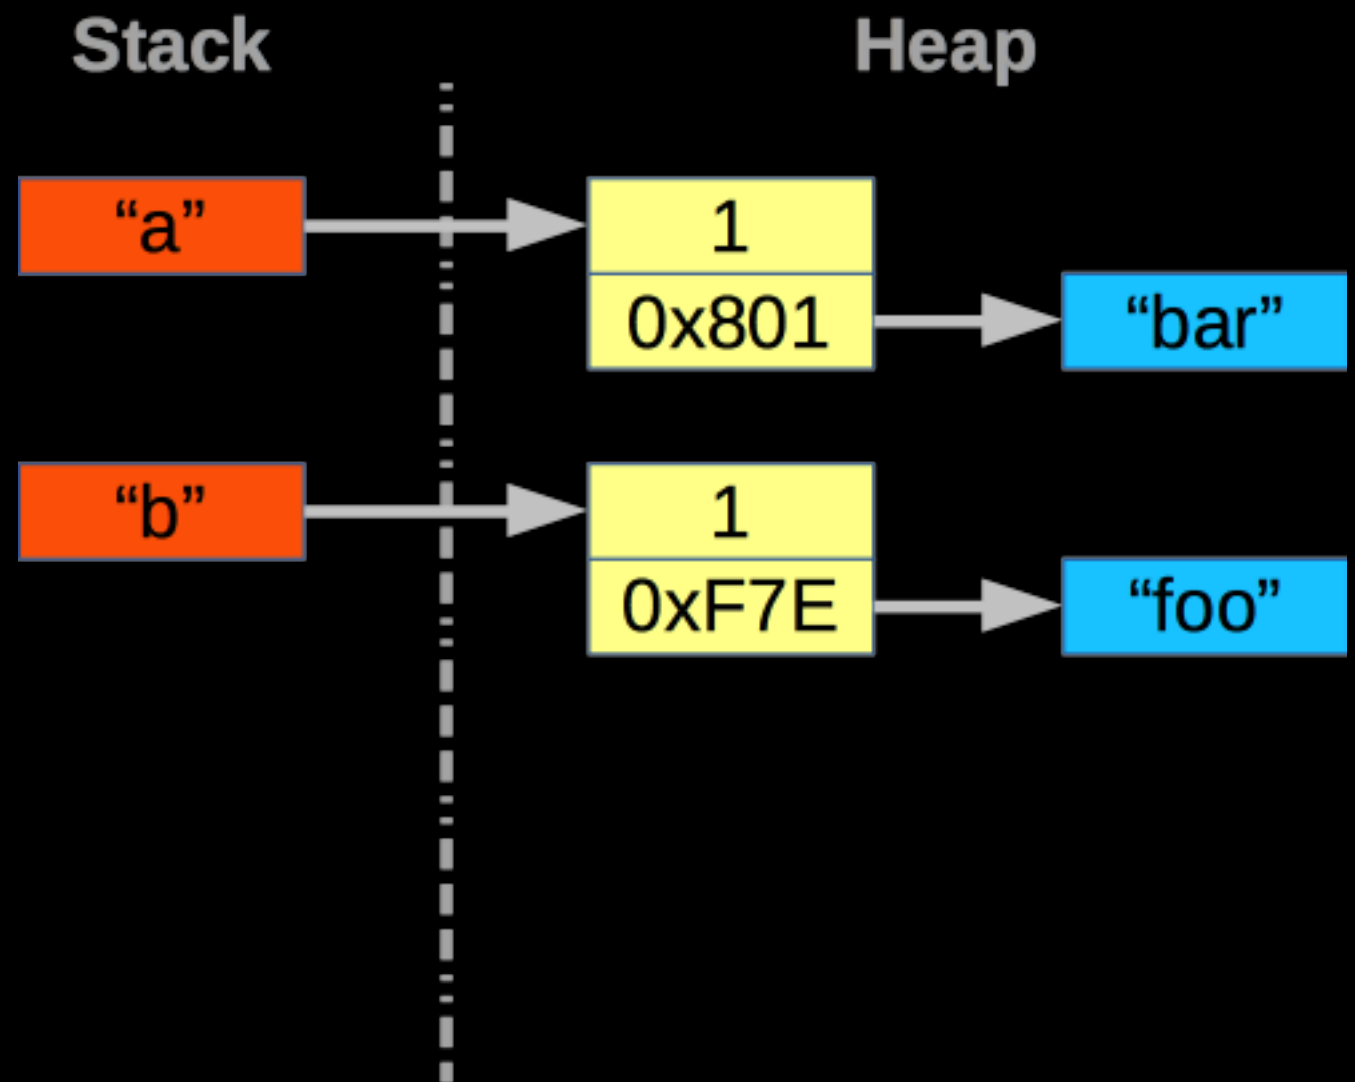
Reference Counting

```
>>> a = 'foo'  
>>> b = a  
>>> a = 'bar'
```



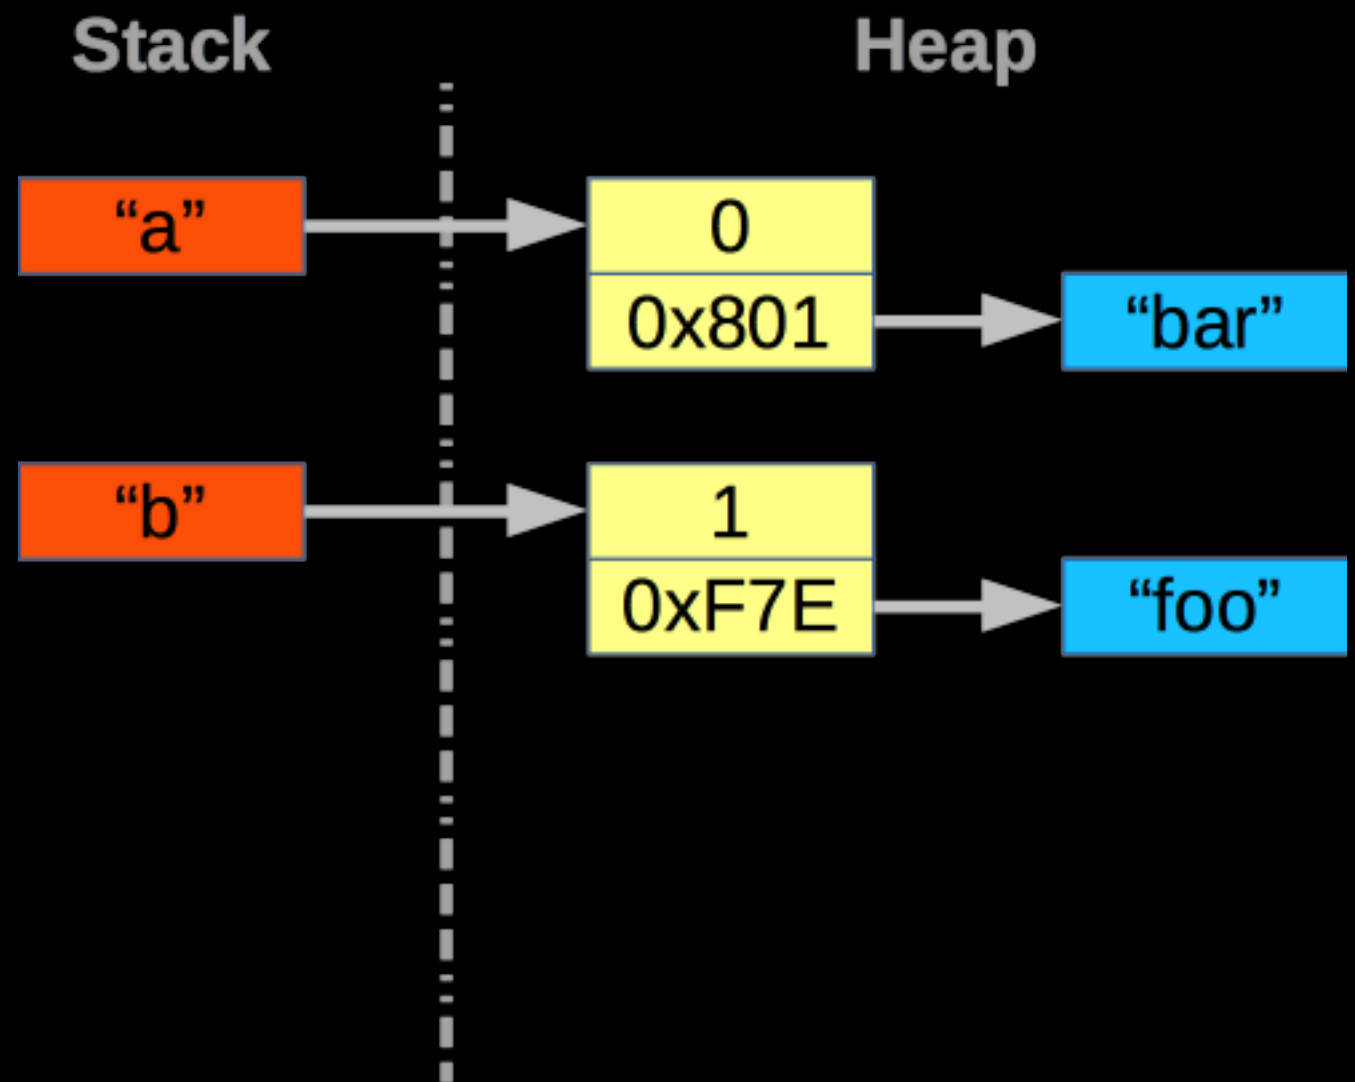
Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
>>> del a
```



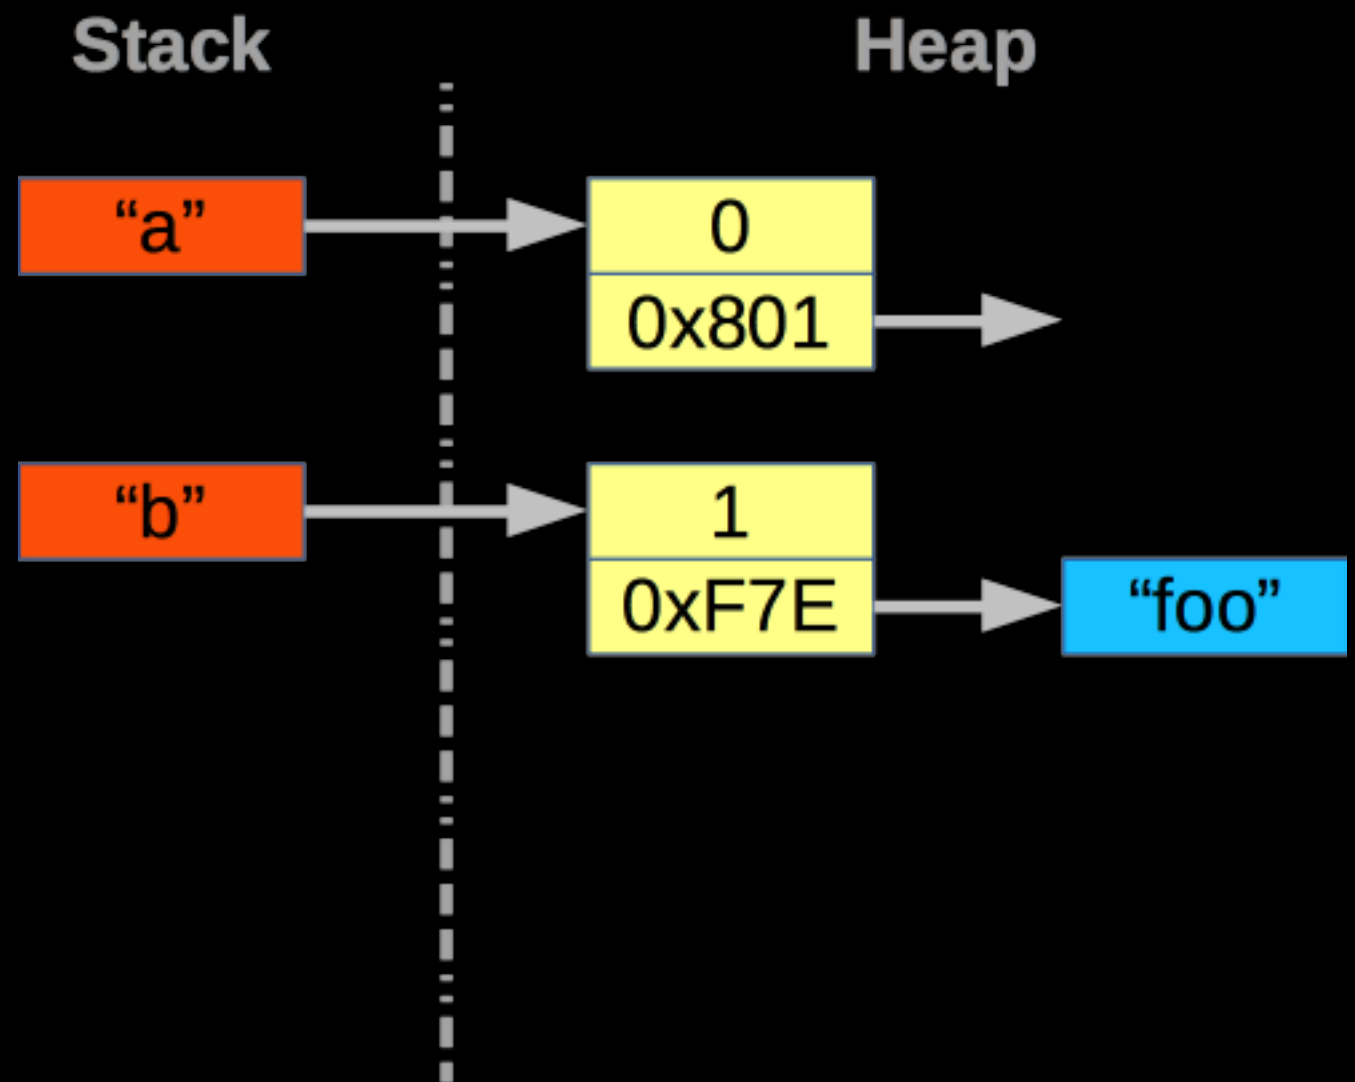
Reference Counting

```
>>> a = 'foo'  
>>> b = a  
>>> a = 'bar'  
>>> del a
```



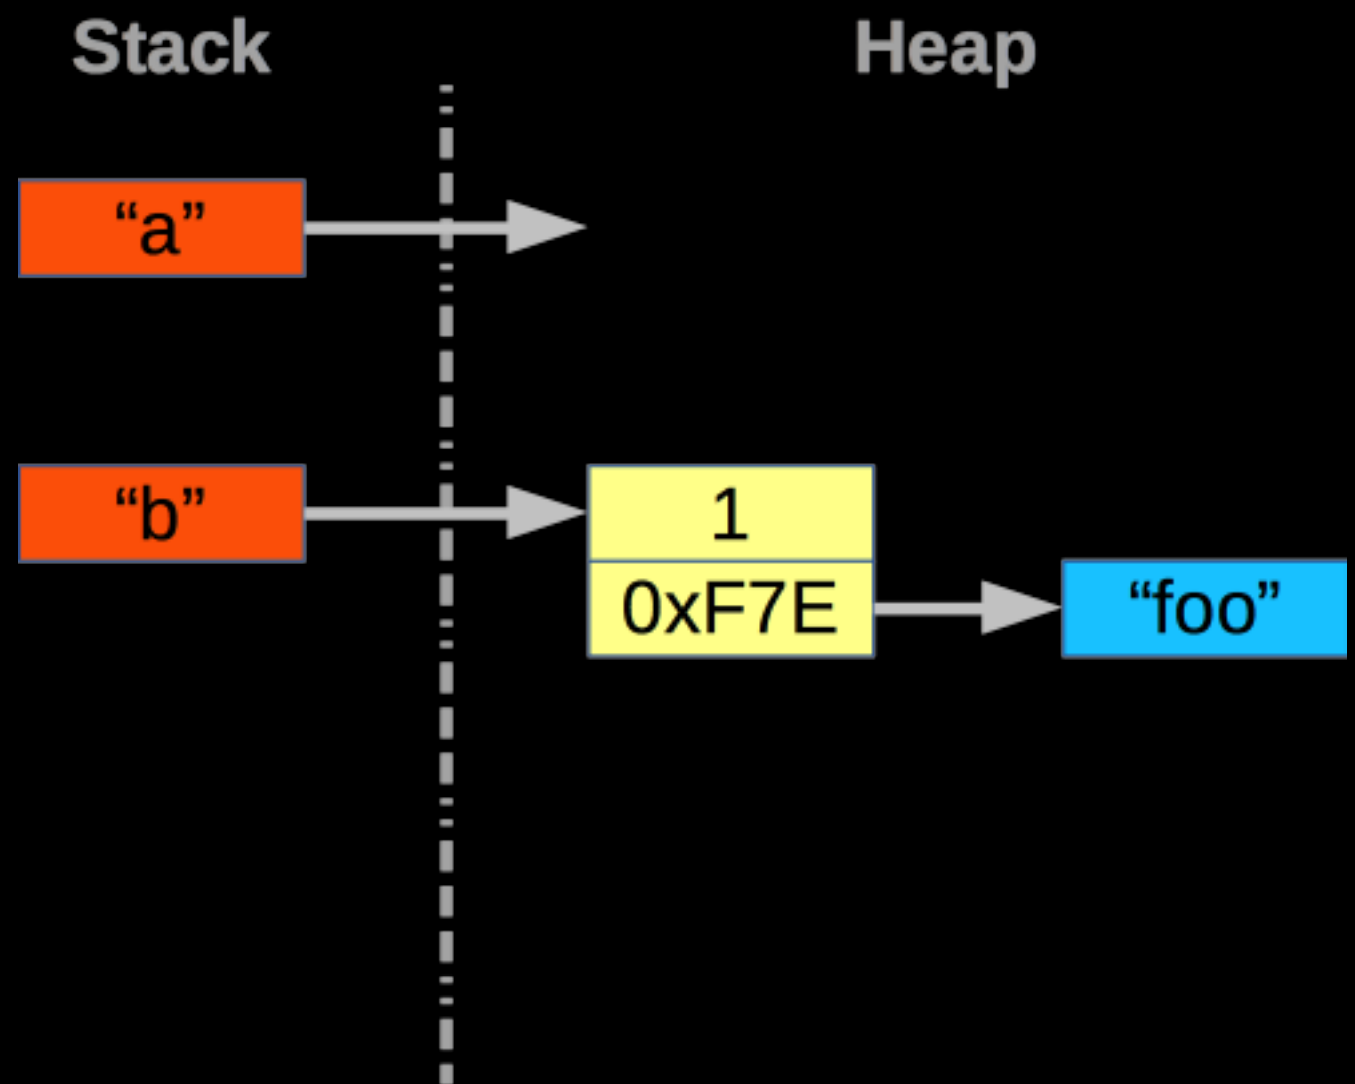
Reference Counting

```
>>> a = 'foo'  
>>> b = a  
>>> a = 'bar'  
>>> del a
```




Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
>>> del a
```



Reference Counting

```
>>> a = 'foo'  
>>> b = a  
>>> a = 'bar'  
>>> del a
```



Garbage Collection

- The GC is just there to resolve cyclic references
 - Only works with containers
- Its not a Unicorn
 - Will not reclaim lost C allocated memory
 - Will not reclaim lost PyObject references

Summary

- 1** coding pattern to keep the dragons at bay
- 2** things to avoid
- 3** kinds of `PyObject*` used in CPython

2 Things to Avoid

- Memory leaks
- Access after `free()`

C Memory Leaks

```
void leak() {  
    char *p = malloc(1024);  
}
```

C Access after free()

```
void access_after_free() {  
    char *p = malloc(1024);  
  
    p[8] = 'A';  
  
    free(p);  
  
    printf("%c", p[8]);  
}
```

Py Memory Leaks

```
#include "Python.h"

void py_leak() {
    PyObject *p0bj;

    p0bj = PyBytes_FromString("Hello world\n");
    PyObject_Print(p0bj, stdout, 0);

}
```

Py Memory Leaks - Fixed

```
#include "Python.h"

void py_leak() {
    PyObject *p0bj;

    p0bj = PyBytes_FromString("Hello world\n");
    PyObject_Print(p0bj, stdout, 0);
    Py_DECREF(p0bj);
}
```

Py Access after Free

```
#include "Python.h"

void py_access_after_free() {
    PyObject *pobj;

    pobj = PyBytes_FromString("Hello world\n");
    PyObject_Print(pobj, stdout, 0);
    Py_DECREF(pobj);
    PyObject_Print(pobj, stdout, 0);
}
```

Py Access after Free

- Don't do this

```
Py_DECREF(p0bj);
```

```
if (p0bj->ob_refcnt > 0) {  
    PyObject_Print(p0bj, stdout, 0);  
}
```

3 Reference Types

- **New** references occur when a `PyObject` is created
 - Example: creating a new list.
- **Stolen** references occur when a `PyObject` is created and assigned. Typically 'setters'
 - Example: setting the value in a list.
- **Borrowed** references are used when getting a `PyObject`
 - Example: accessing a member of a list.
 - If **shared** references mean more to you, great! That's exactly what they are.

Programming by Contract

- **New PyObject*** Your job to free
 - Or give it to someone who will
- **Stolen PyObject*** Who 'steals' it will free it
 - Do not do so yourself
- **Borrowed PyObject*** The real owner can free it at any time
 - Unless you prevent them by registering your interest

New References

```
static PyObject *subtract_long(long a, long b) {  
    PyObject *pA, *pB, *r;  
  
    pA = PyLong_FromLong(a);          /* New reference pA */  
    pB = PyLong_FromLong(b);          /* New reference pB */  
    r = PyNumber_Subtract(pA, pB);     /* New reference r */  
    Py_DECREF(pA);                    /* I have to decref */  
    Py_DECREF(pB);                    /* I have to decref */  
    return r;                          /* Caller has to decref */  
}
```

New References

- Don't do this

```
static PyObject *subtract_long(long a, long b) {  
    return PyNumber_Subtract(PyLong_FromLong(a),  
                               PyLong_FromLong(b));  
}
```

Stolen References

```
static PyObject *make_tuple() {
    PyObject *r, *v;

    r = PyTuple_New(3);          /* New reference r */

    v = PyLong_FromLong(1024L); /* New reference v */
    PyTuple_SetItem(r, 0, v);

    v = PyLong_FromLong(2048L); /* New reference v */
    PyTuple_SetItem(r, 1, v);

    /* More common pattern */
    PyTuple_SetItem(r, 2, PyUnicode_FromString("three"));
    return r;    /* Callers responsibility to decref */
}
```

Stolen References

- Don't do this

```
PyObject *r, *v;
```

```
r = PyTuple_New(3);          /* New reference r */
```

```
v = PyLong_FromLong(1024L); /* New reference v */
```

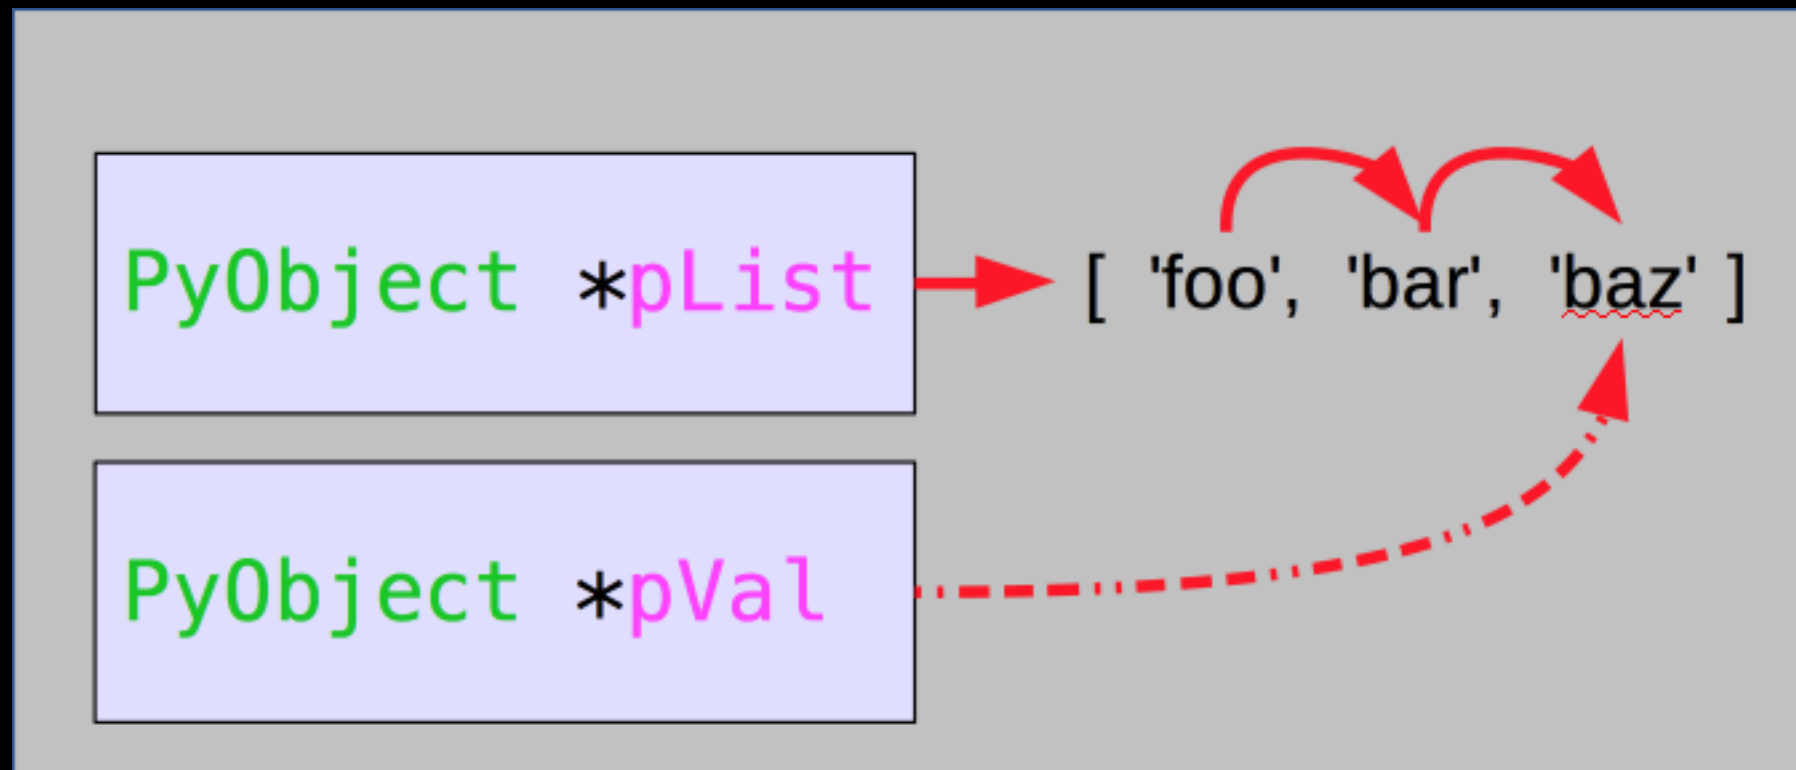
```
PyTuple_SetItem(r, 0, v);    /* r 'steals' the reference */
```

```
Py_DECREF(v);
```

'Borrowed' References

- These are generally 'getters'

```
PyObject *pList = ...  
PyObject *pVal  = PyList_GetItem(pList, 2);
```



'Borrowed' References

- Multiple pointers to the same object - Aaargh!
 - Which is responsible for freeing the object?
 - What happens to the other pointers when one frees the object?
- They can be the source of the most subtle bugs

'Borrowed' References

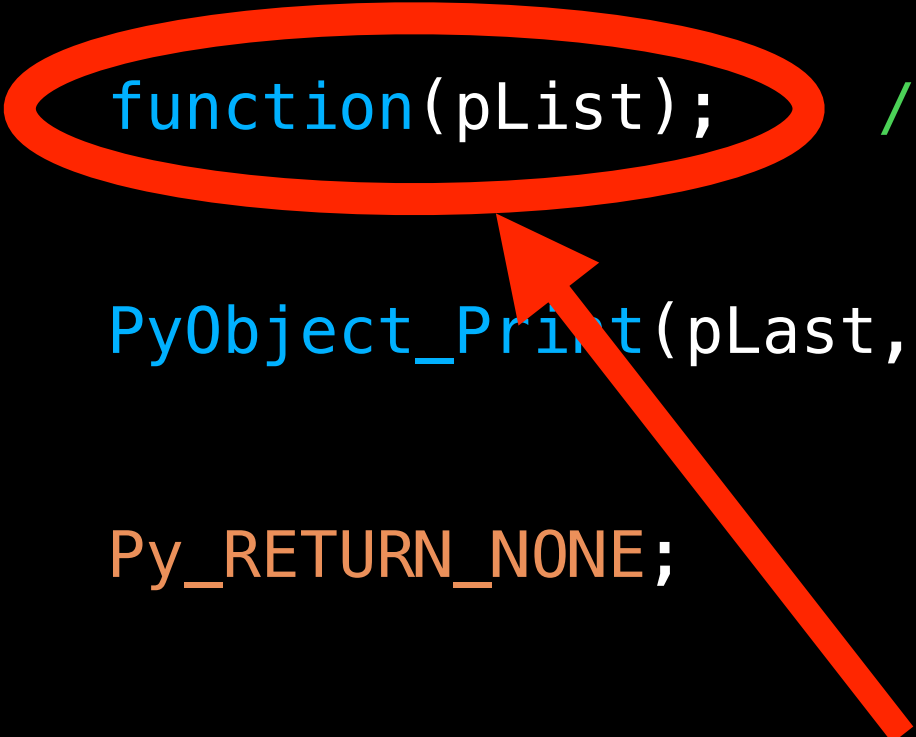
```
static PyObject *borrow_BAD(PyObject *pList) {  
    PyObject *pLast;  
  
    pLast = PyList_GetItem(pList, PyList_Size(pList) - 1);  
  
    function(pList);    /* Dragons ahoy me hearties! */  
  
    PyObject_Print(pLast, stdout, 0);  
  
    Py_RETURN_NONE;  
}
```


Hmm... Suppose

```
static PyObject *borrow_BAD(PyObject *pList) {  
    PyObject *pLast;  
  
    pLast = PyList_GetItem(pList, PyList_Size(pList) - 1);  
  
    function(pList);    /* Dragons ahoy me hearties! */  
  
    PyObject_Print(pLast, stdout, 0);  
  
    Py_RETURN_NONE;  
}
```

Hmm... Suppose

```
static PyObject *borrow_BAD(PyObject *pList) {  
    PyObject *pLast;  
  
    pLast = PyList_GetItem(pList, PyList_Size(pList) - 1);  
  
    function(pList); /* Dragons ahoy me hearties! */  
  
    PyObject_Print(pLast, stdout, 0);  
  
    Py_RETURN_NONE;  
}
```



This deleted the last item in the list

Borrowed Ref Dragon 0

```
>>> import cPyRefs  
>>> l = ['foo', 'bar', 'baz']  
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

Borrowed Ref Dragon 1

```
>>> import cPyRefs
>>> l = ['foo', 'bar', 'baz']
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

```
>>> import cPyRefs
>>> l = ['foo', 'bar', 'baz']
>>> a = l[-1]
>>> cPyRefs.borrow_bad(l) # Works fine!
```

Borrowed Ref Dragon 2

```
>>> import cPyRefs  
>>> l = ['foo', 'bar', 'baz']  
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

Borrowed Ref Dragon 2

```
>>> import cPyRefs  
>>> l = ['foo', 'bar', 'baz']  
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

```
>>> import cPyRefs  
>>> l = [1, 2, 3]  
>>> cPyRefs.borrow_bad(l) # Works fine!
```

Borrowed Ref Dragon 2

```
>>> import cPyRefs
>>> l = ['foo', 'bar', 'baz']
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

```
>>> import cPyRefs
>>> l = [1, 2, 3]
>>> cPyRefs.borrow_bad(l) # Works fine
```

```
>>> import cPyRefs
>>> l = [800, 801, 802]
>>> cPyRefs.borrow_bad(l) # Kaboom!
```

Run-time Errors + Data Dependent Errors

Run-time Errors + Data Dependent Errors



The Problem

```
static PyObject *borrow_BAD(PyObject *pList) {  
    PyObject *pLast;  
  
    pLast = PyList_GetItem(pList, PyList_Size(pList) - 1);  
  
    function(pList);    /* Dragons ahoy me hearties! */  
  
    PyObject_Print(pLast, stdout, 0);  
  
    Py_RETURN_NONE;  
}
```

The Fix

```
static PyObject *borrow_BAD(PyObject *pList) {
    PyObject *pLast;

    pLast = PyList_GetItem(pList, PyList_Size(pList) - 1);
>    Py_INCREF(pLast);    /* Protect myself */

    function(pList);    /* Dragons tamed! */

    PyObject_Print(pLast, stdout, 0);

>    Py_DECREF(pLast);    /* I'm done */
>    pLast = NULL;

    Py_RETURN_NONE;
}
```

1 Pattern For Reliable C

- Borrowed references incref'd and decref'd correctly.
- A single place for clean up code
 - No early returns
- Exception consistency. Either:
 - An exception is set **and** NULL is returned.
 - Or: no Exception set **and** non-NULL returned.

Writing Pythonic Python

```
def simple(obj):  
    ret = None;  
  
    try:  
        # Do fabulous stuff here  
        # On error raise an exception  
    except ... as err:  
        # Handle exceptions  
    finally:  
        # And we are out  
    return ret;
```

Writing Pythonic C

Writing Pythonic C

```
static PyObject *simple(PyObject *arg1) {  
    PyObject *ret = NULL;  
  
    goto try;  
try:  
    /* Do fabulous stuff here */  
    /* On error "goto except;" */  
    goto finally;  
except:  
    /* Handle exceptions */  
finally:  
    /* And we are out. */  
    return ret;  
}
```

Function Entry

```
static PyObject *func(PyObject *arg1) {  
    /* Create any local PyObject* as NULL */  
    PyObject *obj_a      = NULL;  
    /* Create the PyObject* return value as NULL */  
    PyObject *ret        = NULL;  
  
    goto try; /* Pythonic 'C' ;-) */  
try:
```


try:

try:

```
assert(! PyErr_Occurred());
/* Inc the reference count of the arguments. */
assert(arg1);
Py_INCREF(arg1);

/* Your code here */

/* Local object creation; borrowed or new. */
obj_a = ...;
/* If an error ... */
if (! obj_a) {
    PyErr_SetString(PyExc_ValueError, "0oops.");
    goto except;
}
```

try:

```
/* Return object creation, ret will either be a
 * new reference or a borrowed reference
 * INCREf'd */
ret = ...;
if (! ret) {
    PyErr_SetString(PyExc_ValueError,
                    "0oops again.");
    goto except;
}
/* If success then check exception is clear,
 * goto finally; with non-NULL return value. */
assert(! PyErr_Occurred());
assert(ret);
goto finally;
except:
```

except:

except:

```
/* Failure so Py_XDECREF the return value */  
Py_XDECREF(ret);  
/* Check a Python error set somewhere above */  
assert(PyErr_Occurred());  
/* Signal failure */  
ret = NULL;  
/* Fall through to finally: */
```

finally:

finally:

finally:

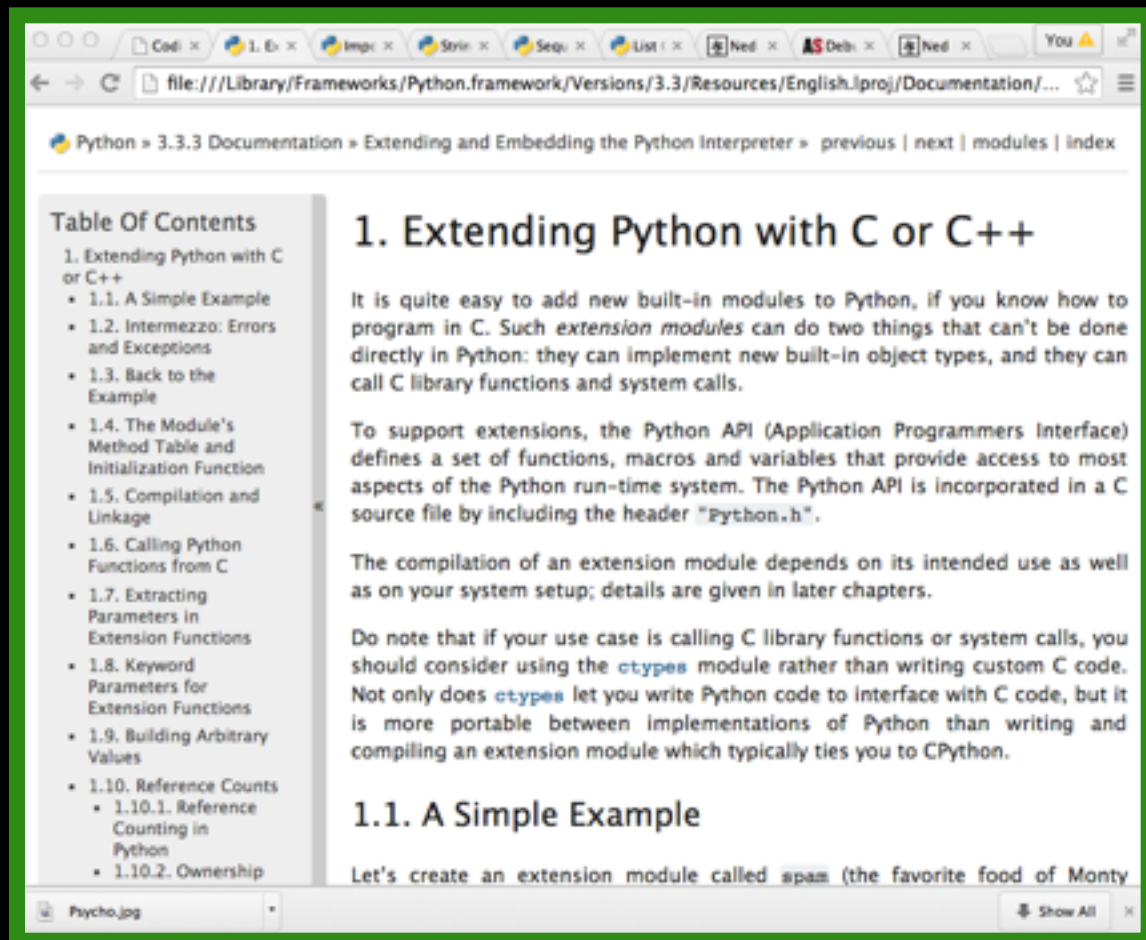
```
/* All _local_ PyObject's are Py_XDECREF'd here.
 * For new references this will free them.
 * For borrowed references this
 * will return them to their previous state. */
Py_XDECREF(obj_a);
/* Decrement the ref count of given arguments
 * if they have been incremented. */
Py_DECREF(arg1);
/* And return... */
return ret;
}
```

All this and more...

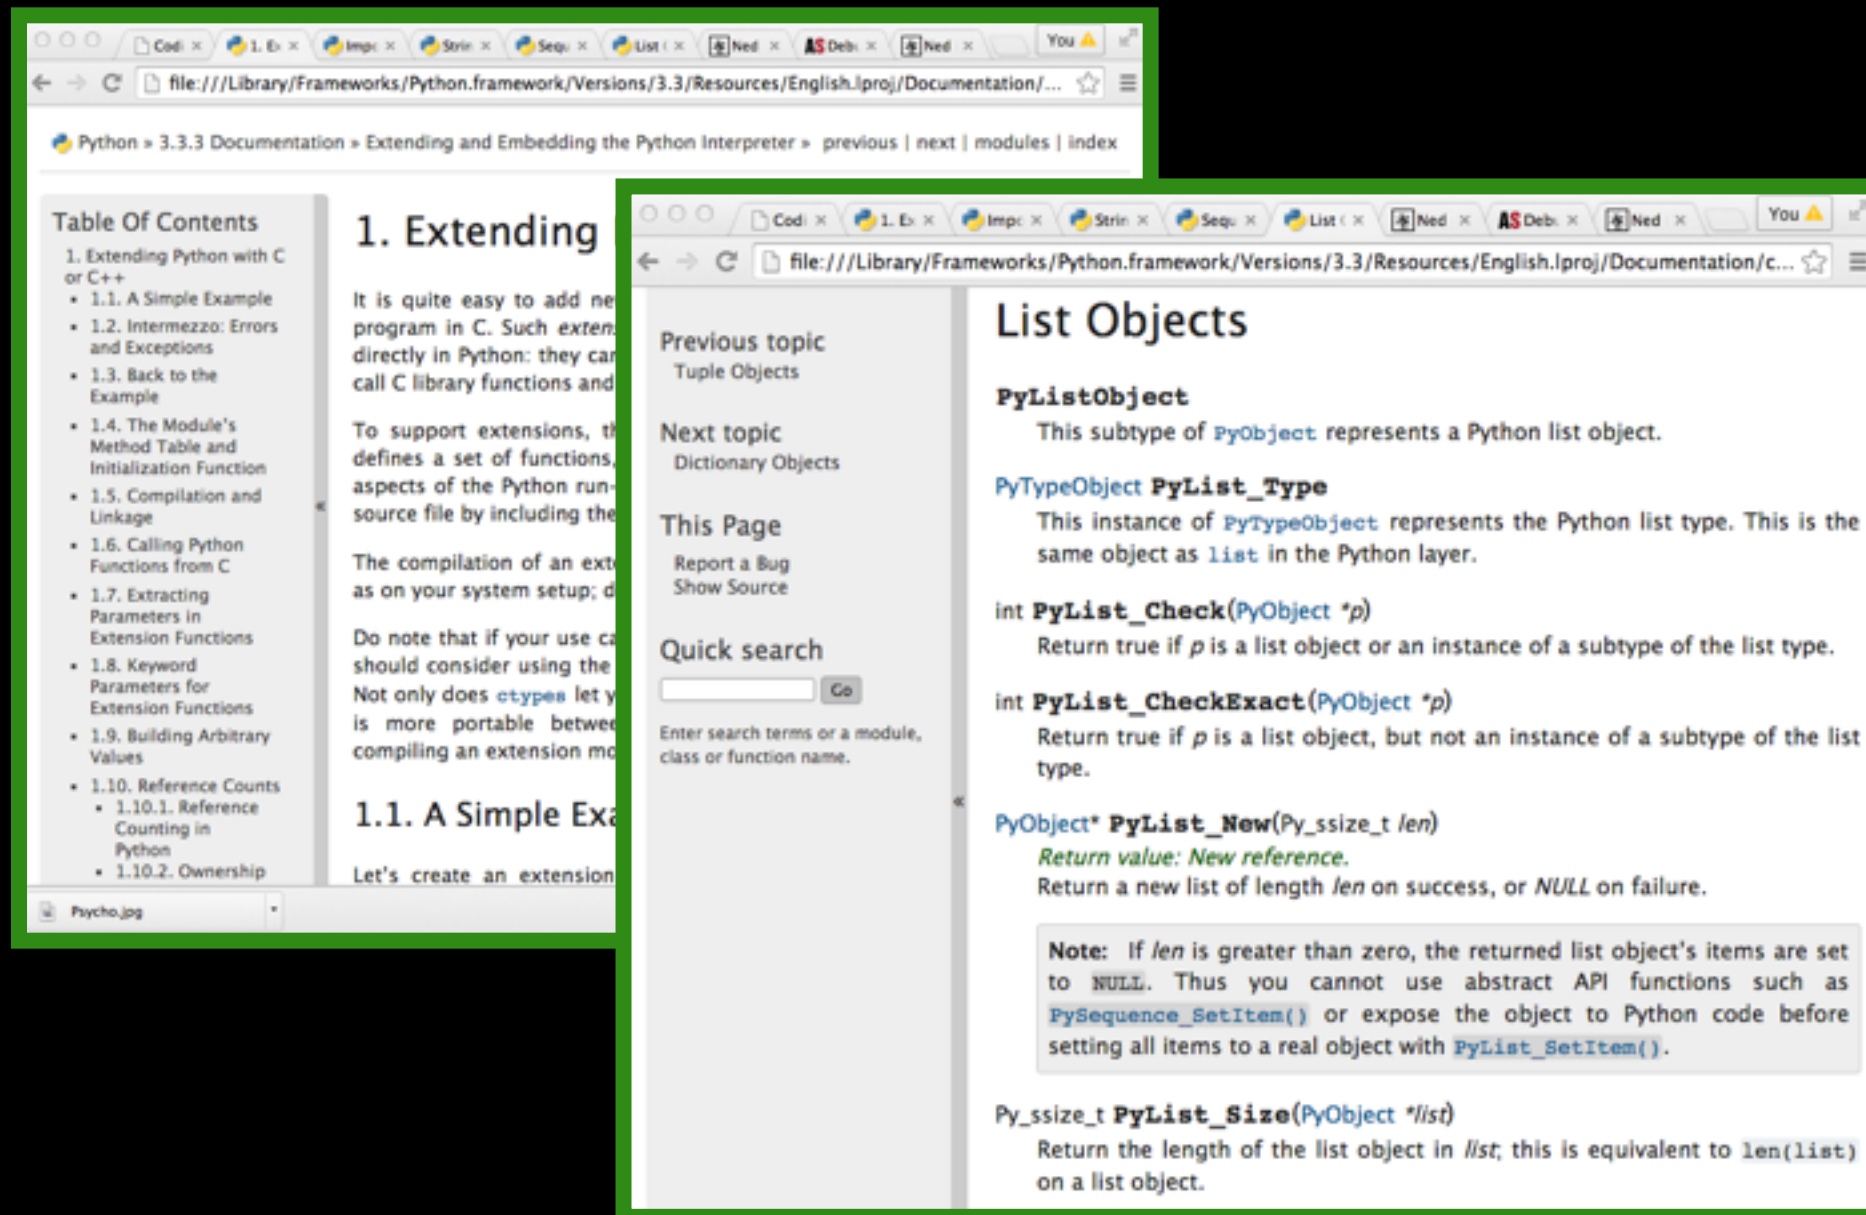
<https://github.com/paulross>

In “PythonExtensionPatterns”

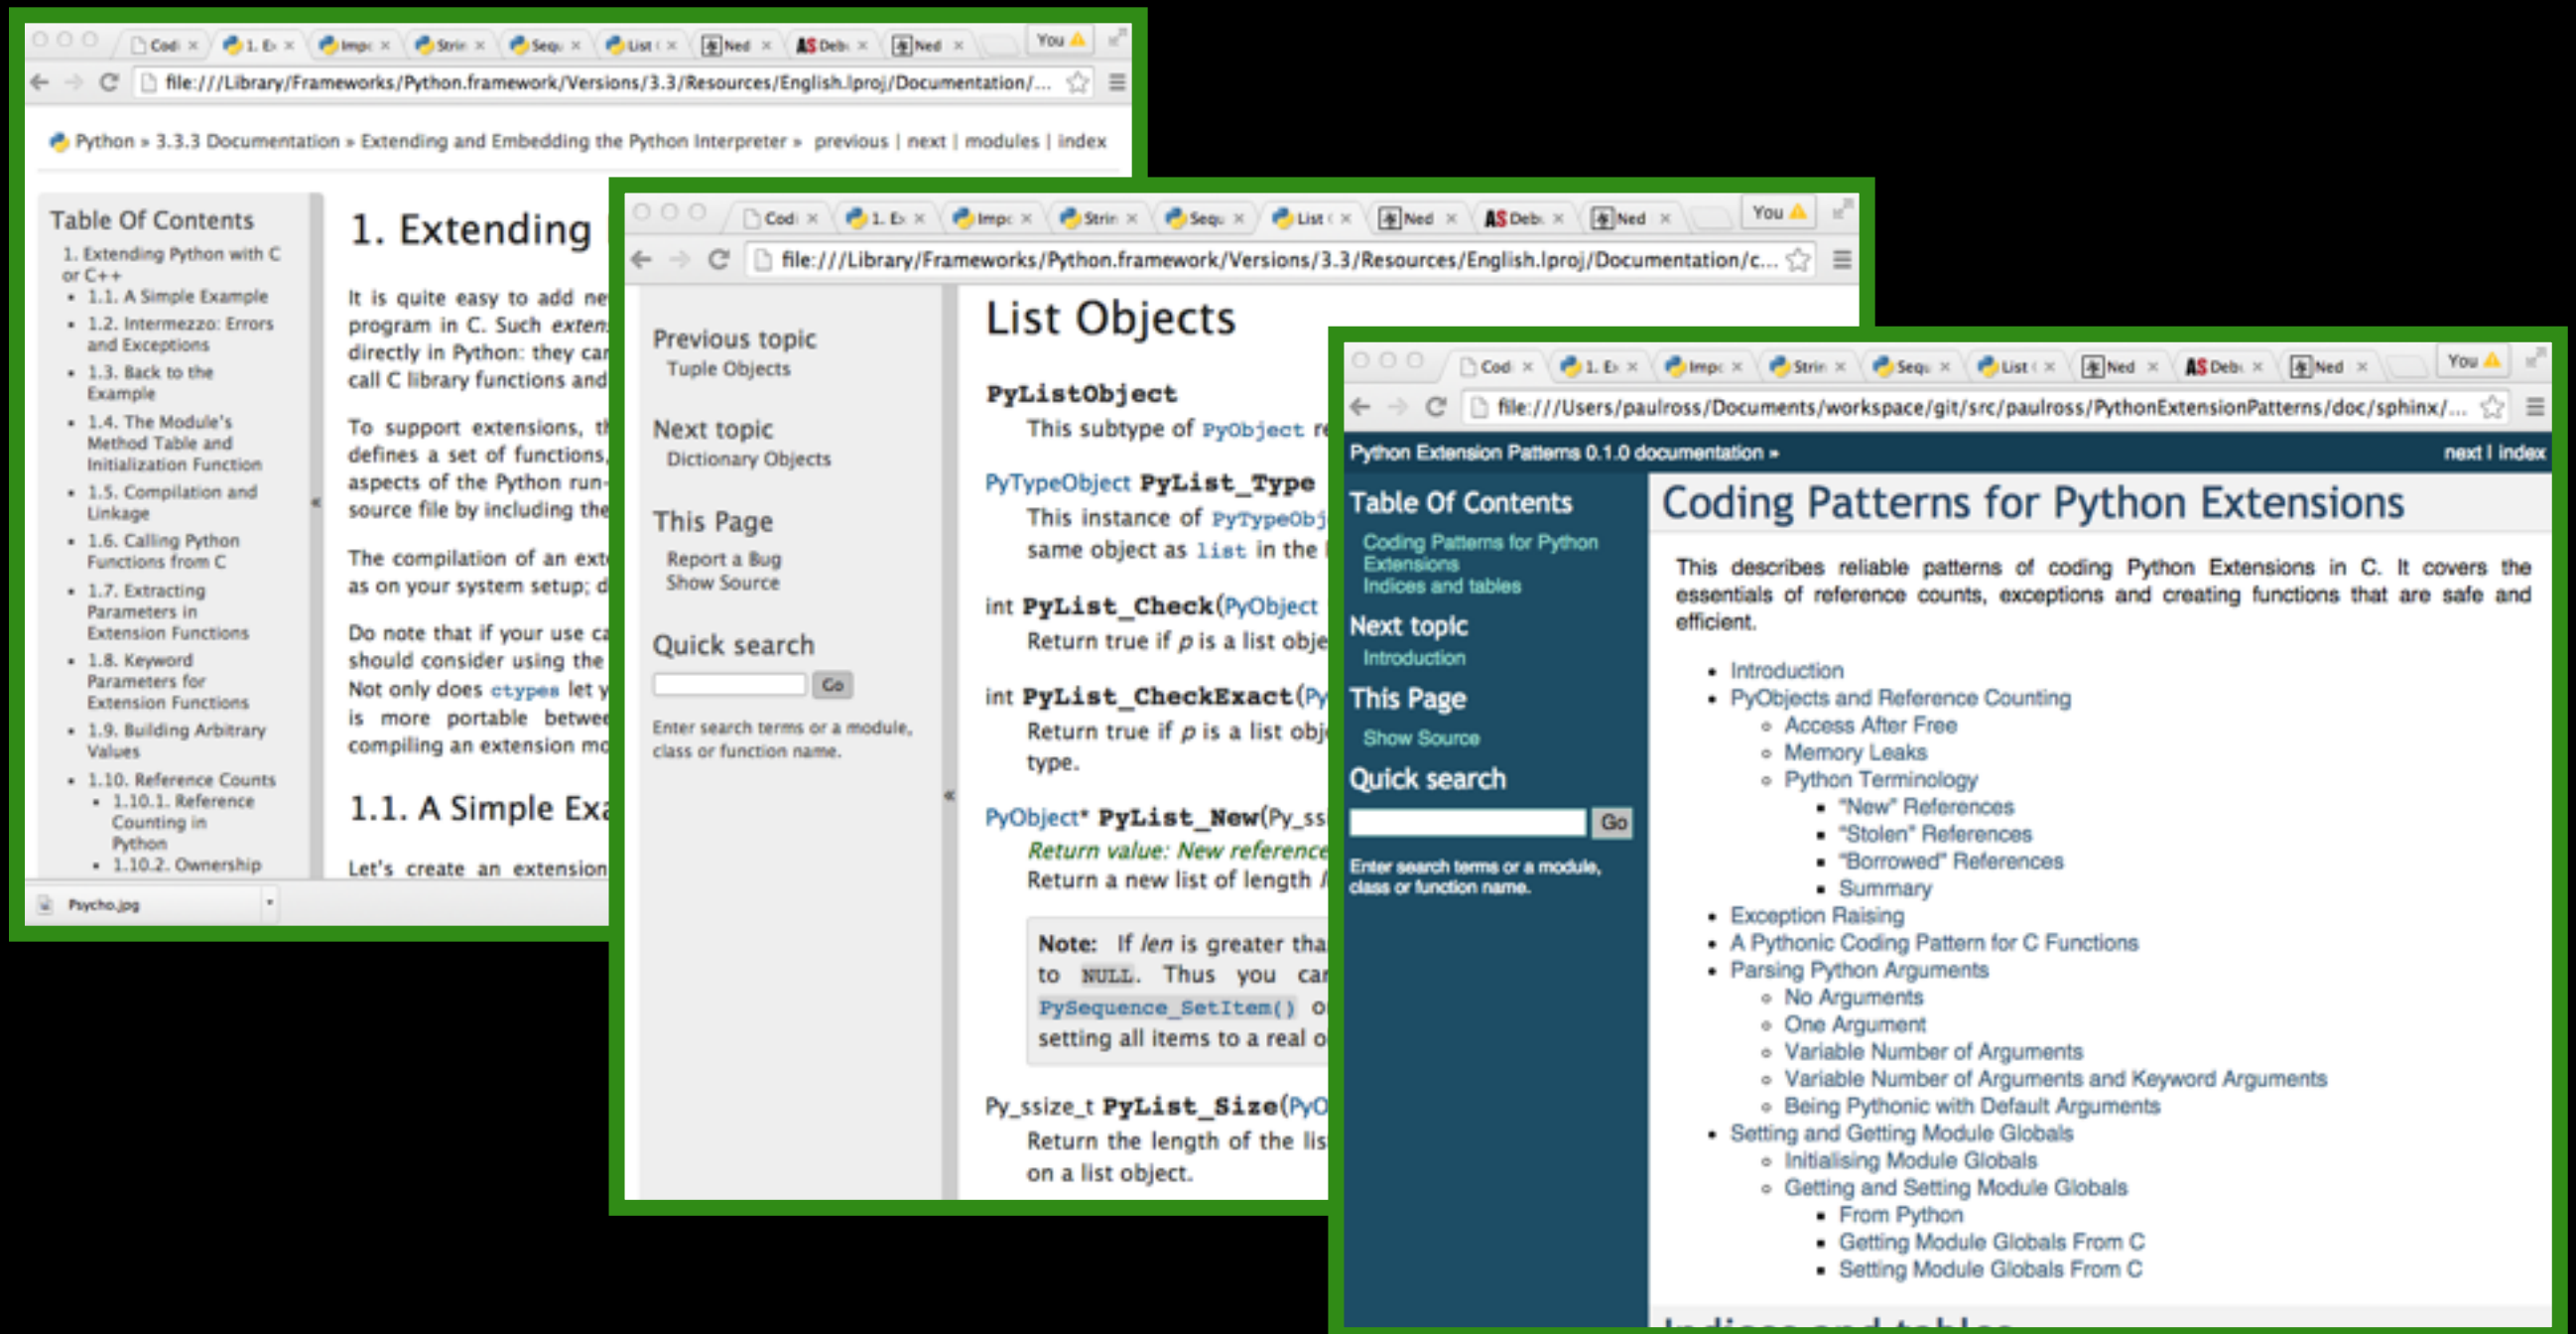
The Documentation is Excellent - Use it!



The Documentation is Excellent - Use it!




The Documentation is Excellent - Use it!



War Story ~ Mandatory

- Proprietary binary files of oilfield data
 - Self describing, variable format, sequentially written
- Make them random access by creating an index
 - The index is built with a sequence of `seek()`/`read()` operations
 - `read()` is about 1% to 2% of the original file size
- Originally written in Python. Typ. 10-100ms/Mb
- How fast can we go?

Improvement in indexing cost using C extension cPhysRec.



Summary

1 coding pattern to keep the dragons at bay

2 things to avoid

- `malloc()` with no `free()`
- Access after `free()`

3 kinds of references to `PyObject*`

- **New**: its yours
- **Stolen**: its theirs
- **Borrowed**: you are sharing something that is really theirs - let them know!

That's It



<https://github.com/paulross/PythonExtensionPatterns>