Here Be Dragons

Writing Reliable Python Extensions in C

Paul Ross, Lead Technologist, AHL

PyCon 2016 Portland, Oregon

Man AHL

https://twitter.com/manahltech

- London based systematic hedge fund since 1987
- $19.2bn Funds Under Management (2016-03-31)
- We are active in 400+ markets in 40+ countries
- We take ~2bn market data points each day
  - https://github.com/manahl/arctic
- 125 people, 22 first languages. And Python!

Why Write in C?

- Blinding performance
- Interface with C/C++ libraries
- Leaner resources
- Flee the GIL

Here Be Dragons

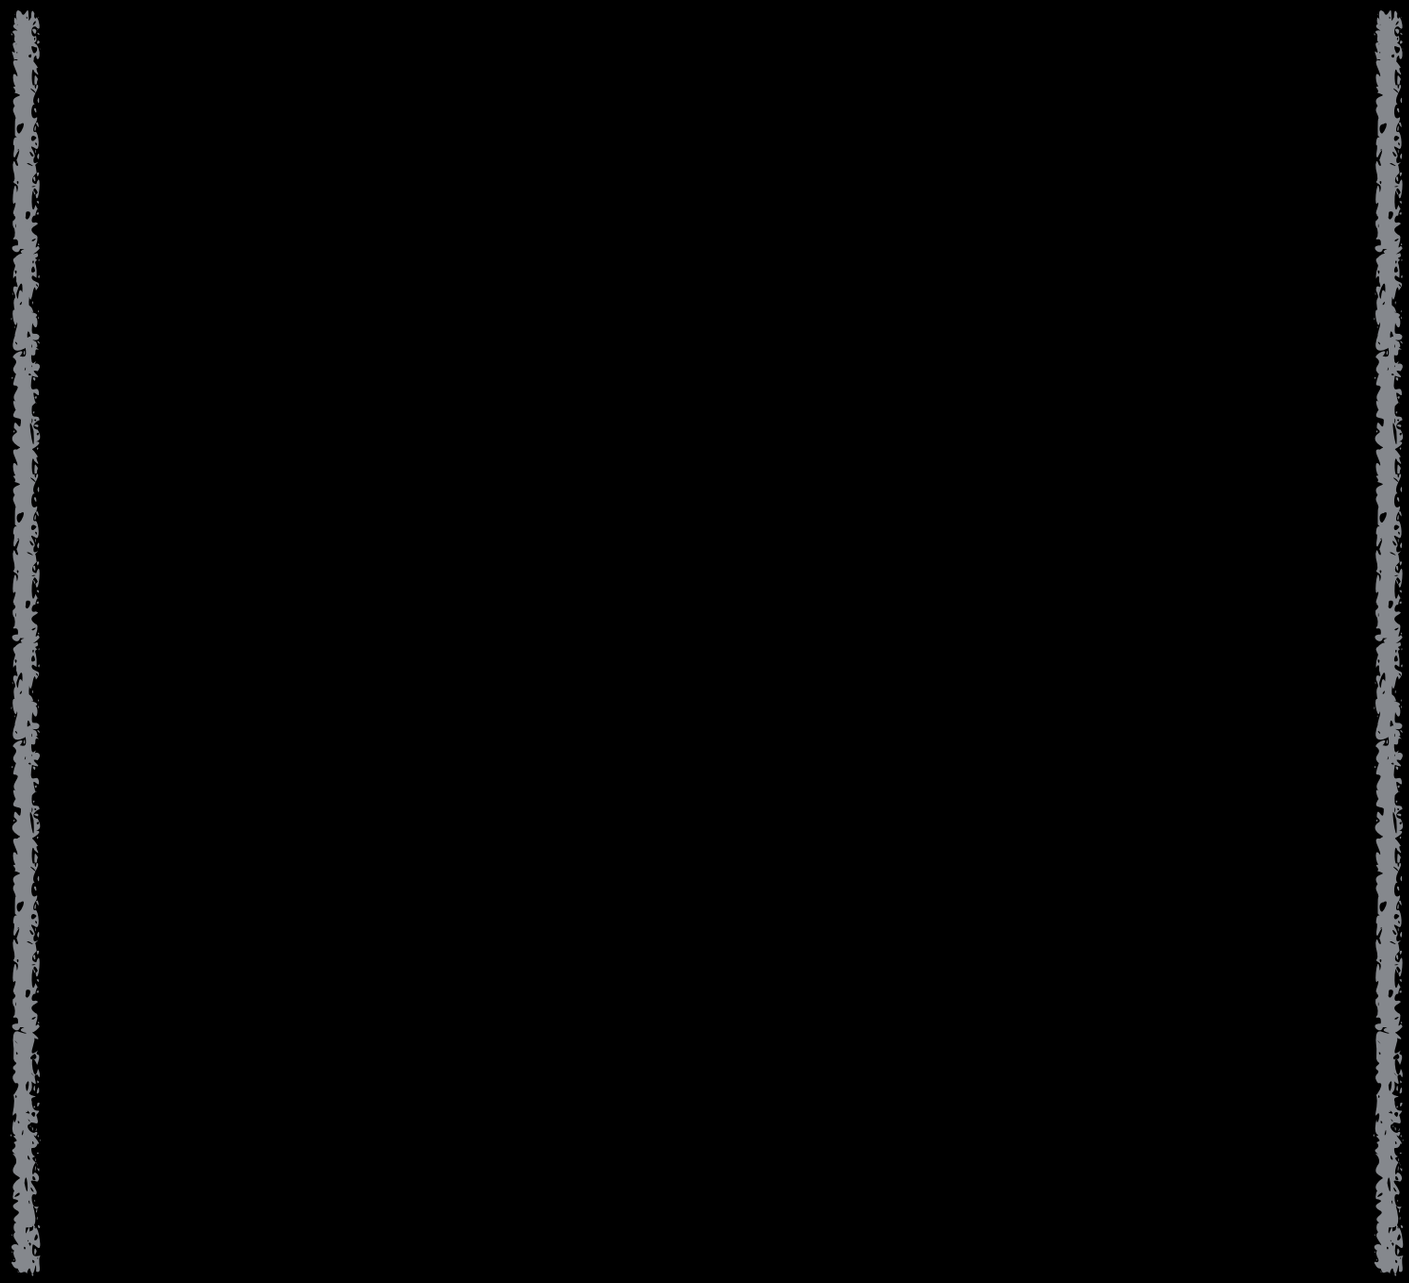

Here Be Dragons

Here Be Dragons

C

Here Be Dragons

CPython code!

C

Here Be Dragons

CPython code!

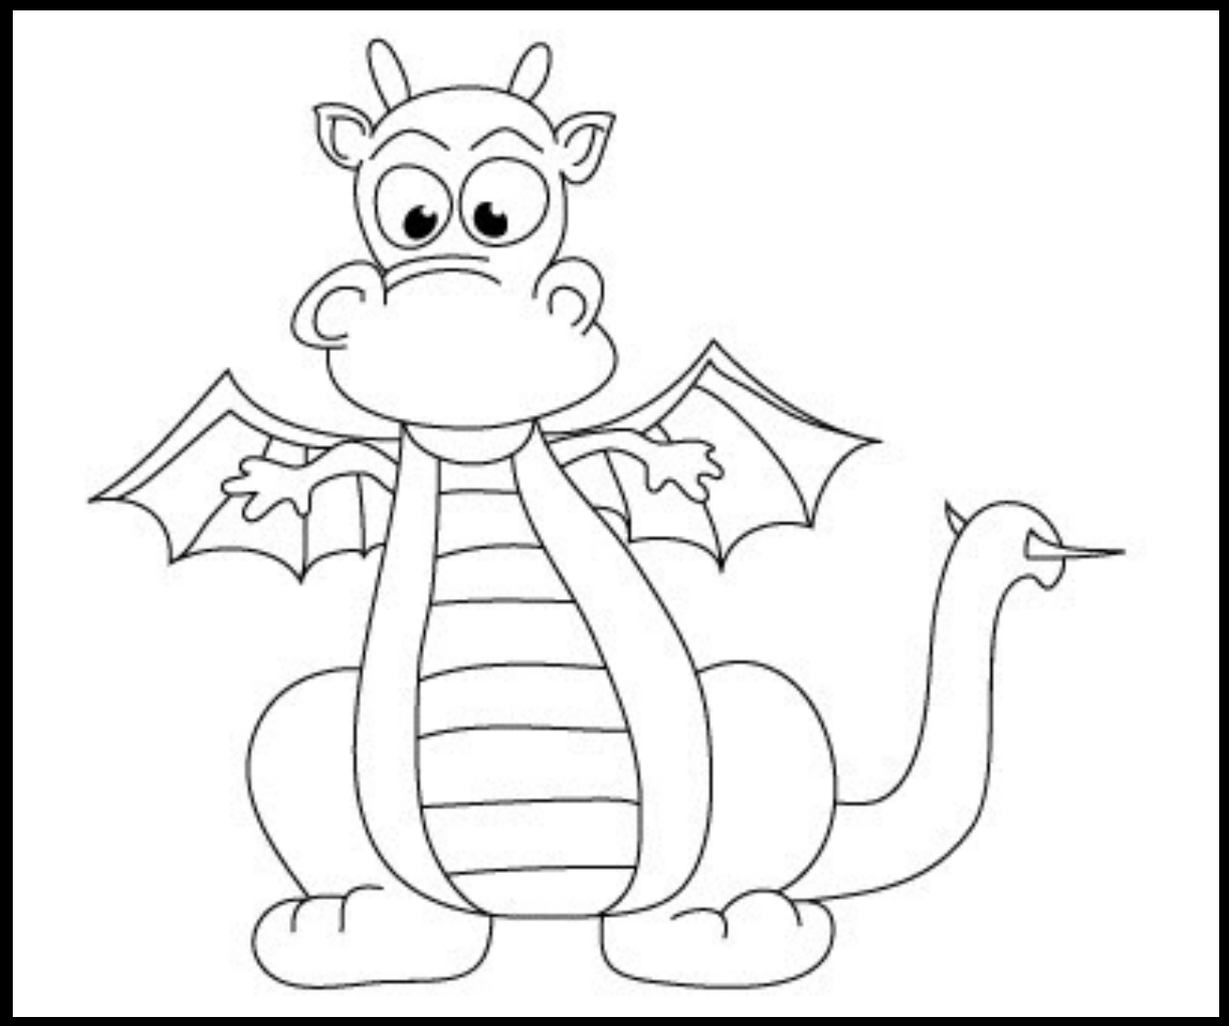

Automatic Memory Management

- Every Python object has a reference count
  - On creation this is set to 1
  - When it becomes 0 the object can be deallocated

Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
>>> del a
```

Reference Counting

```
>>> a = 'foo'
```

Reference Counting

```
>>> a = 'foo'
```

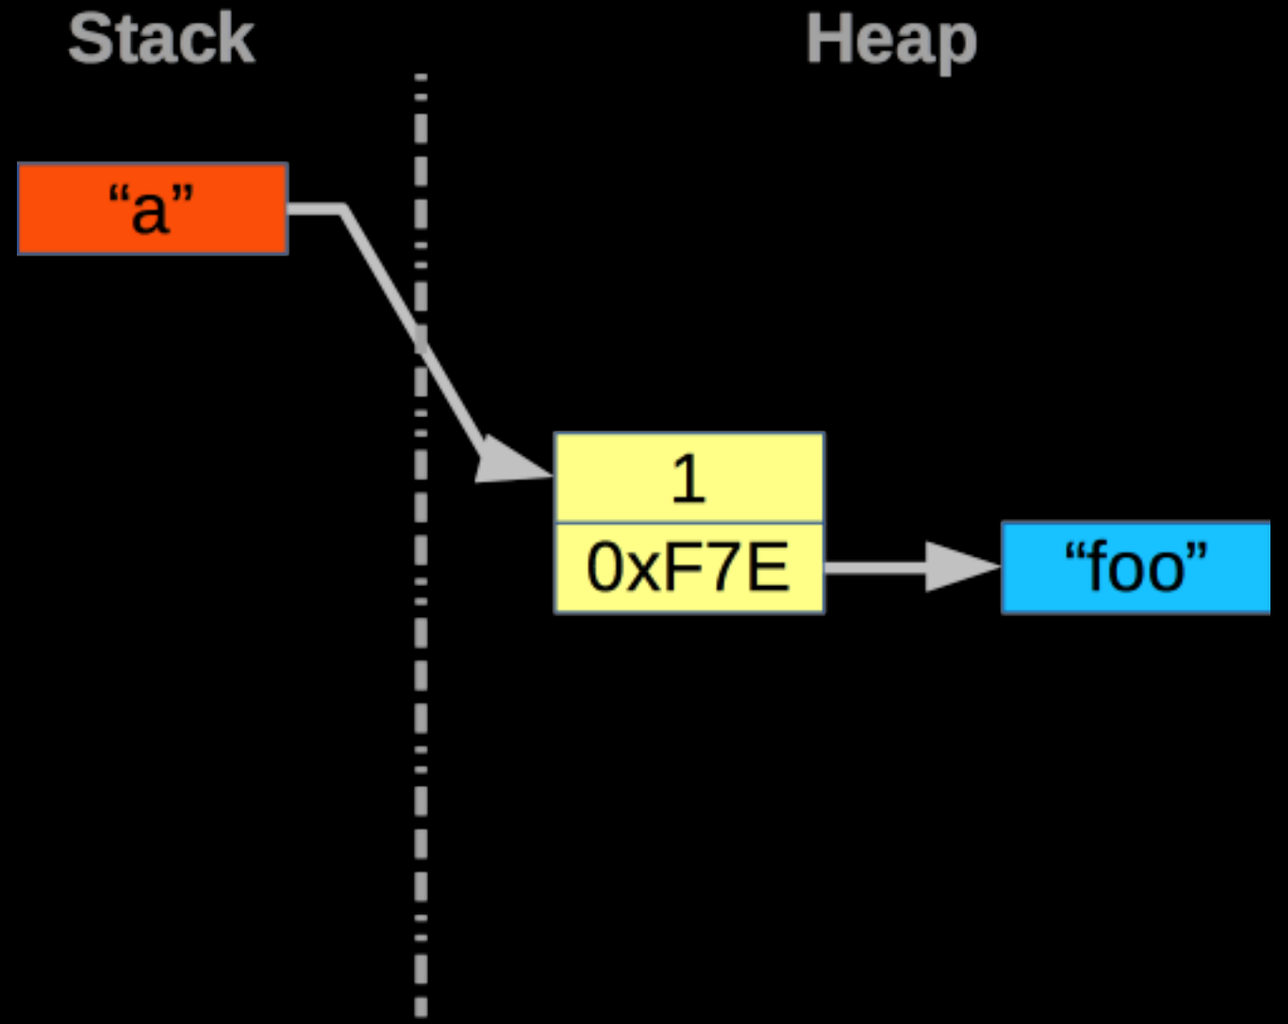

Reference Counting

```
>>> a = 'foo'
```

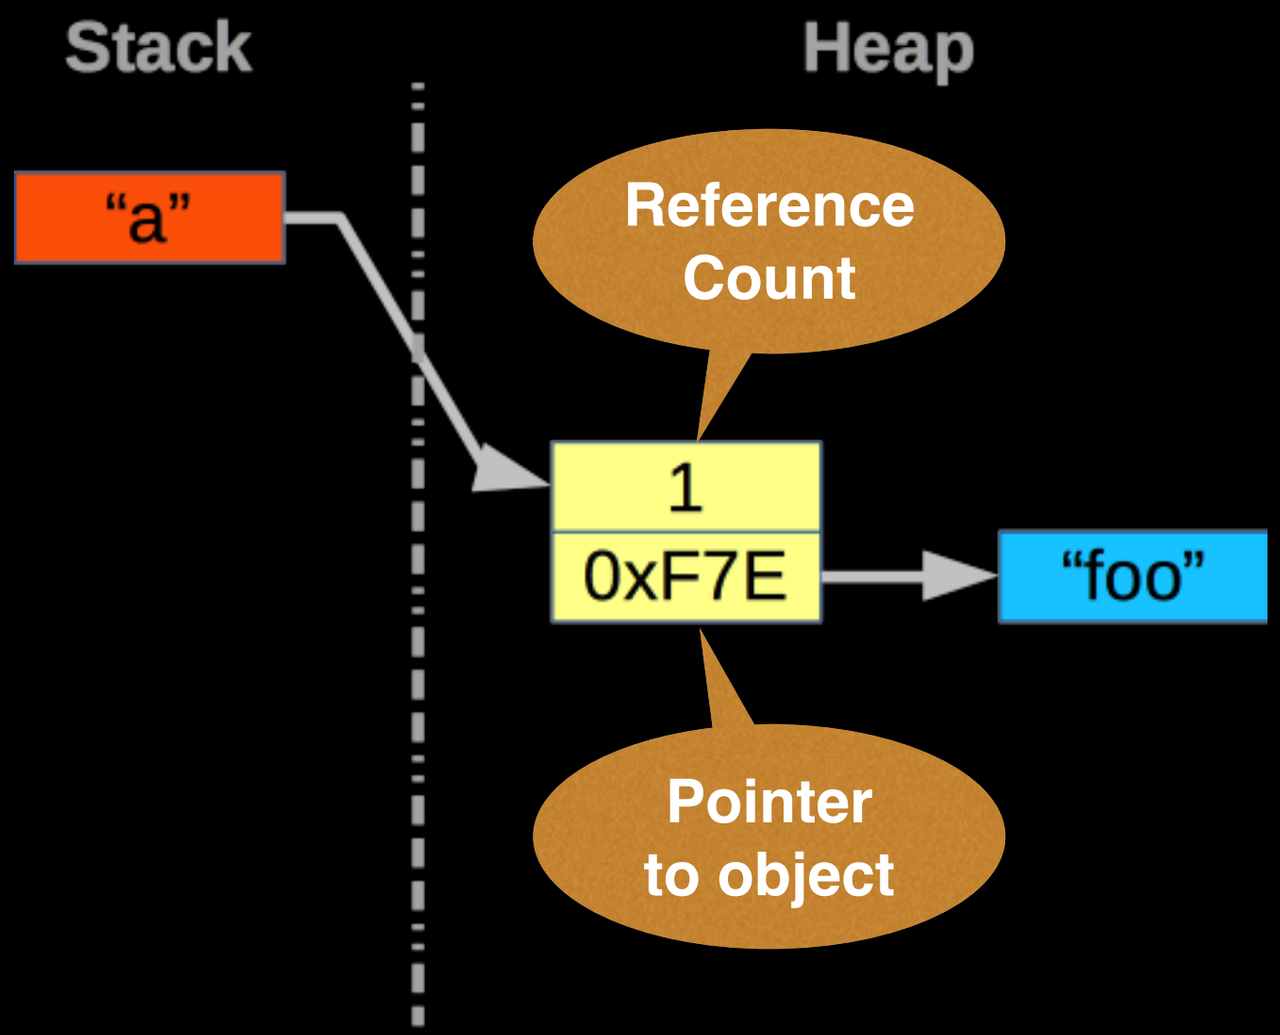

Reference Counting

```
>>> a = 'foo'
>>> b = a
```

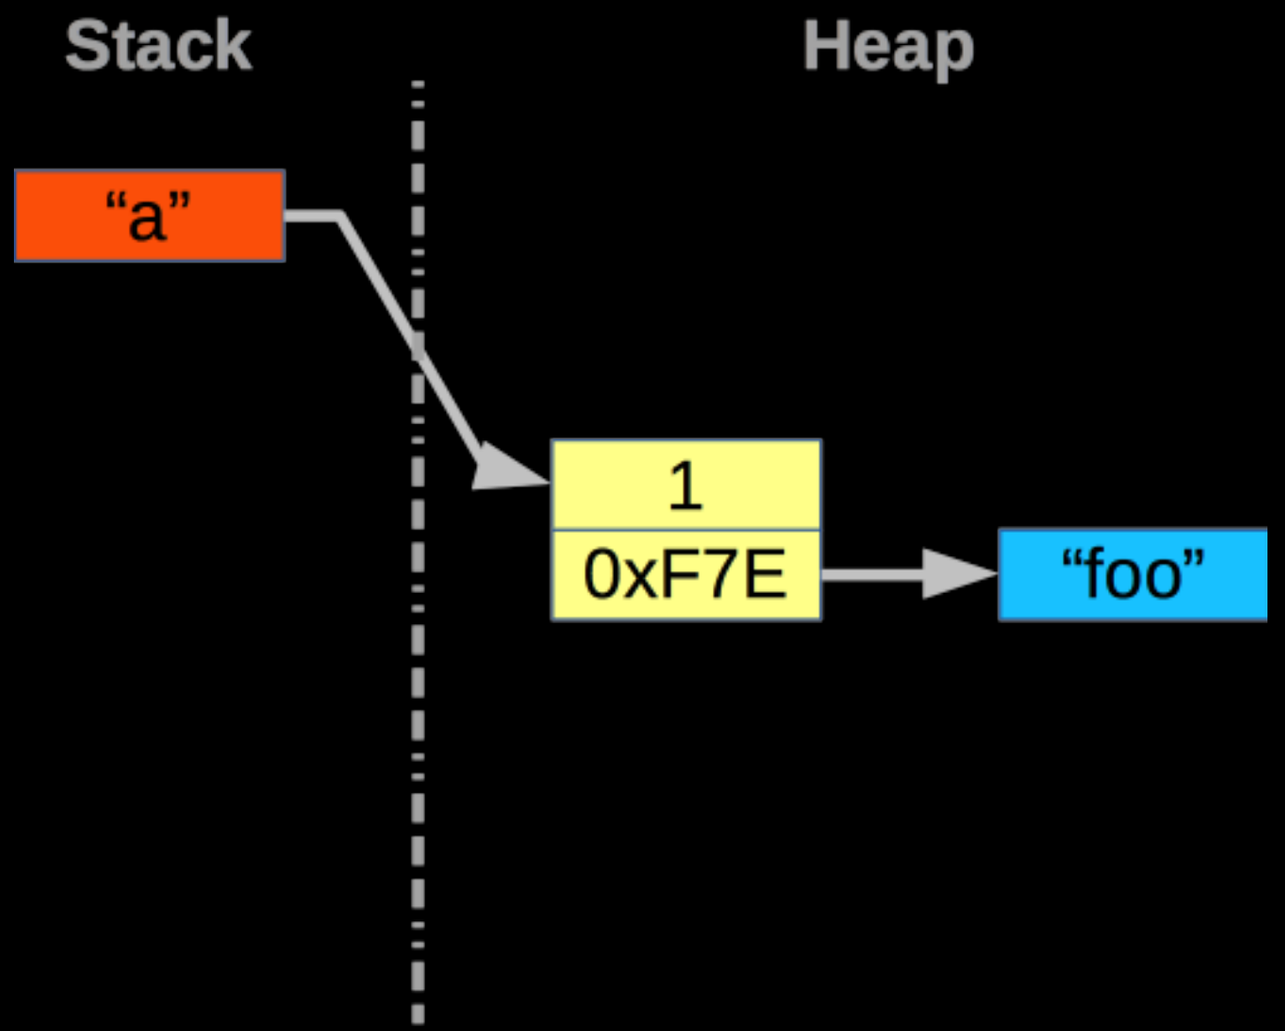

Reference Counting

```
>>> a = 'foo'
>>> b = a
```

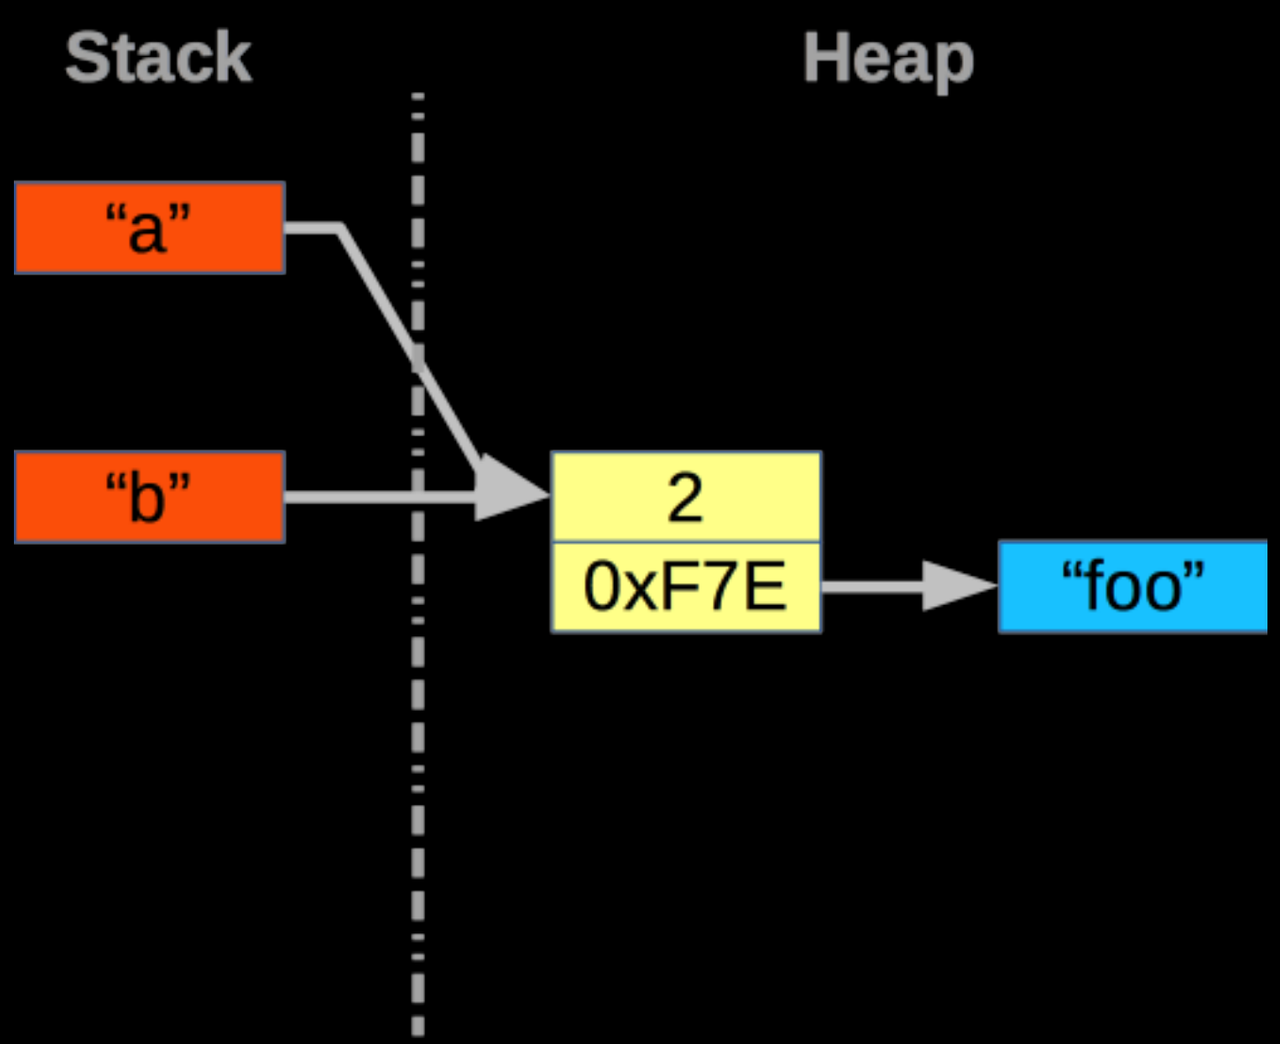

Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
```

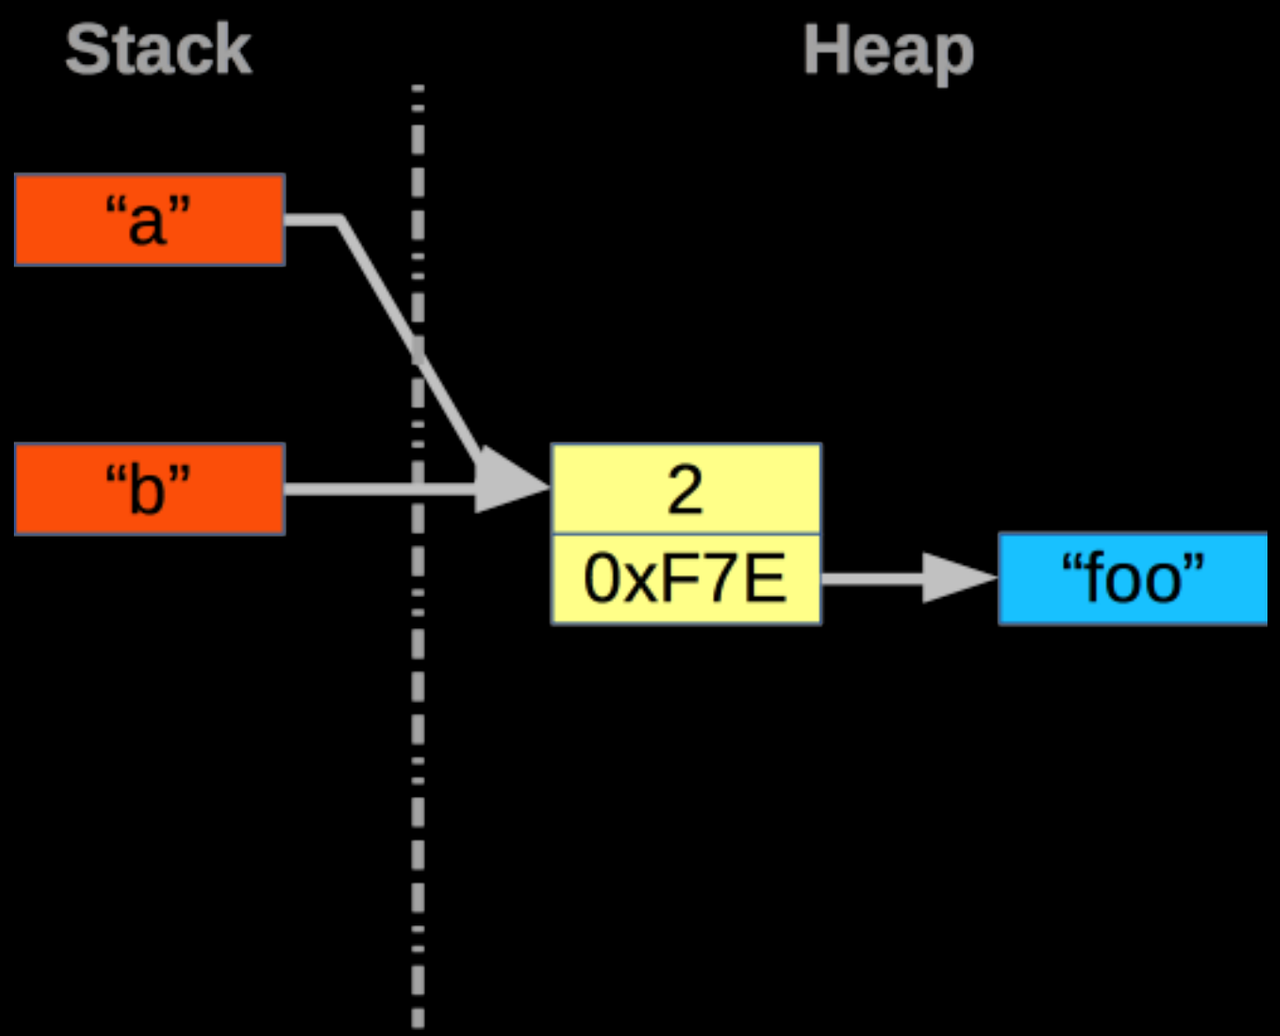

Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
```

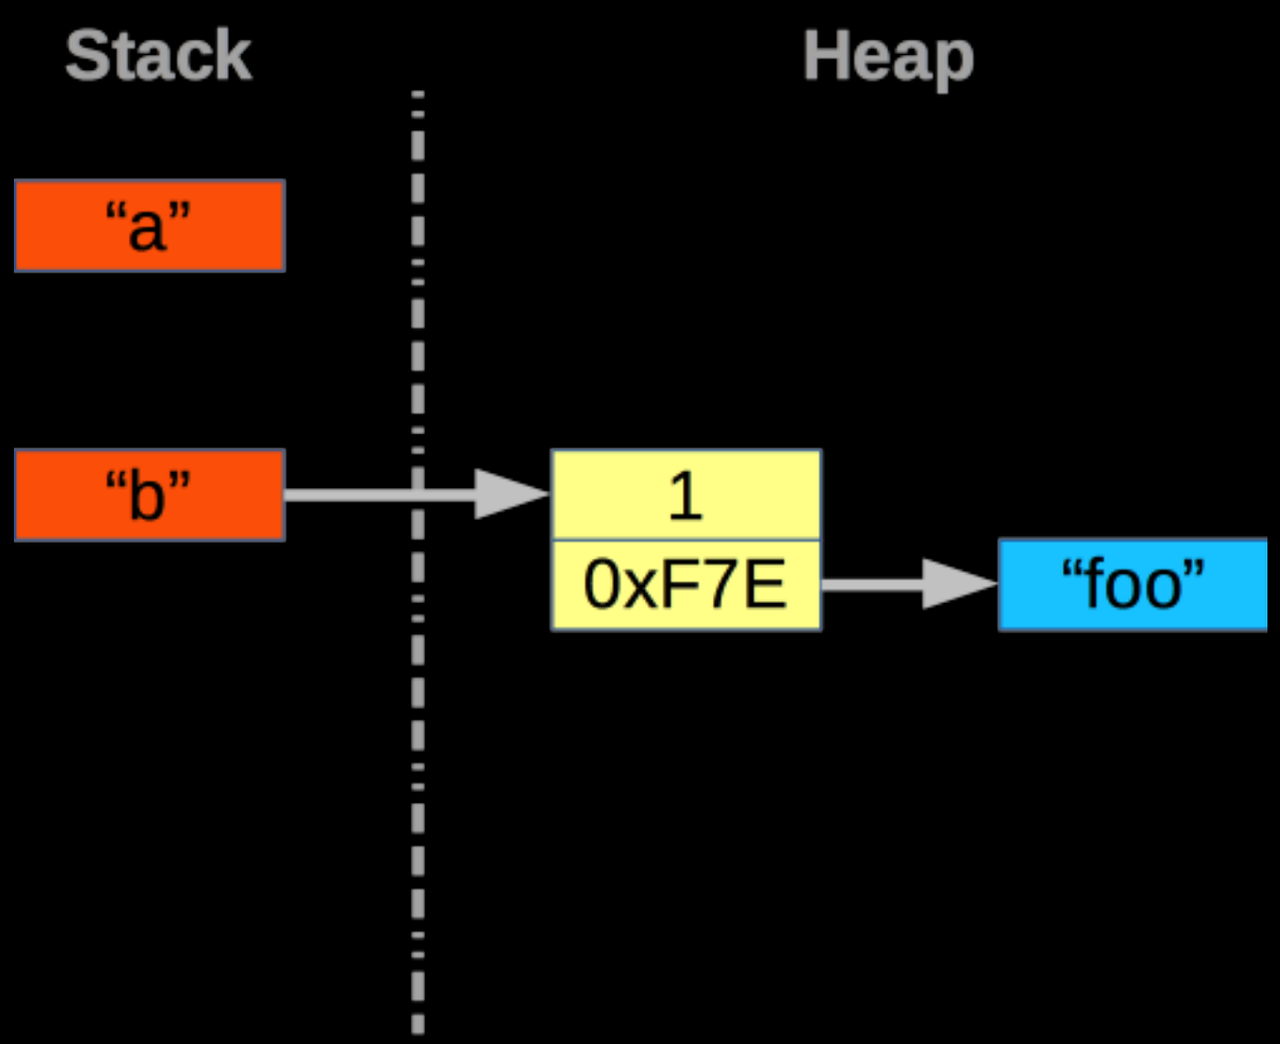

Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
```

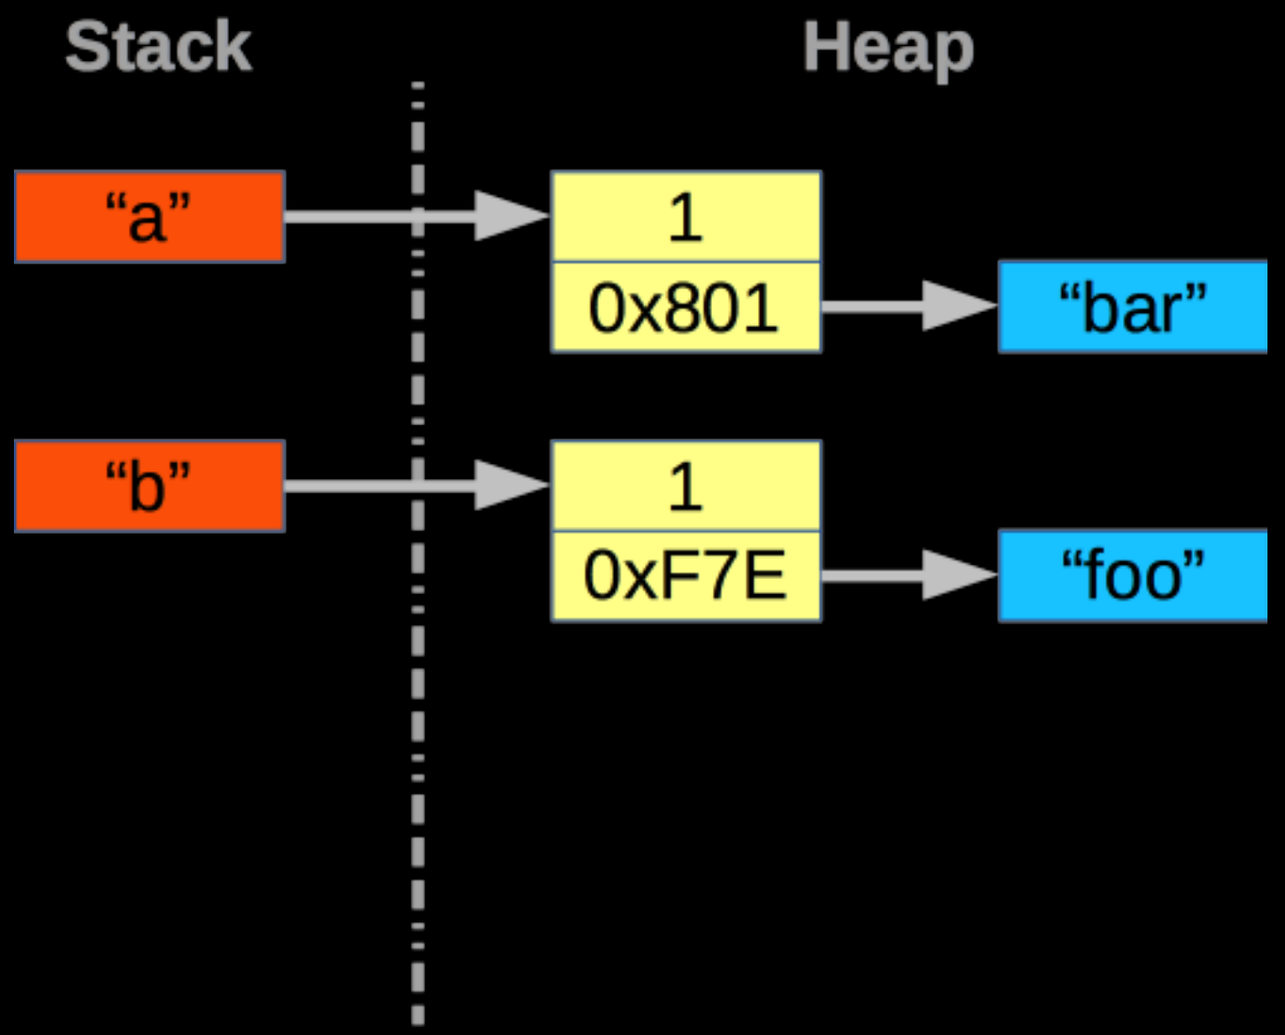

Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
>>> del a
```

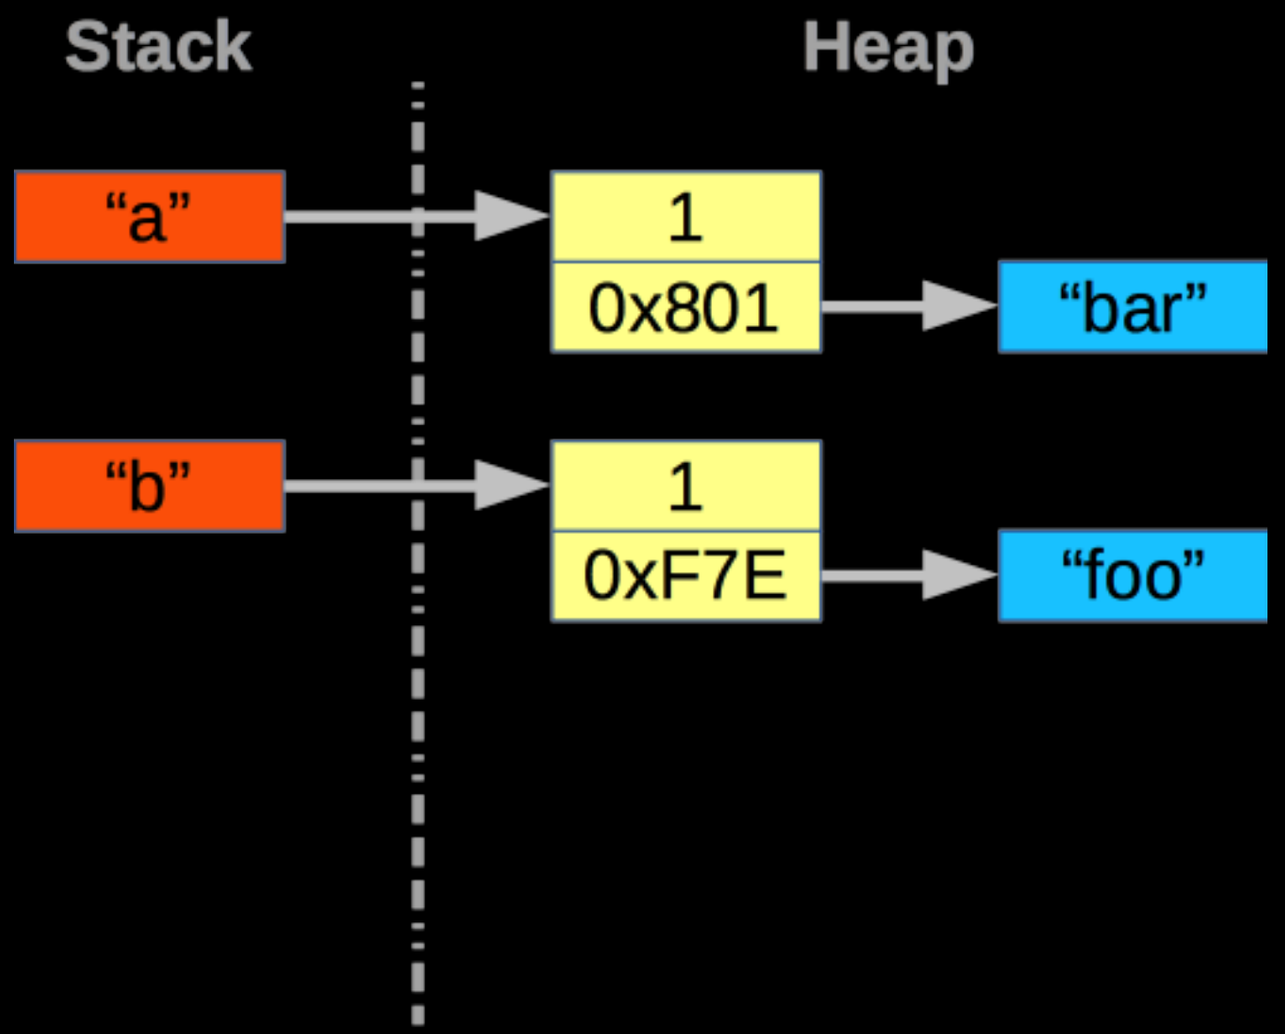

Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
>>> del a
```

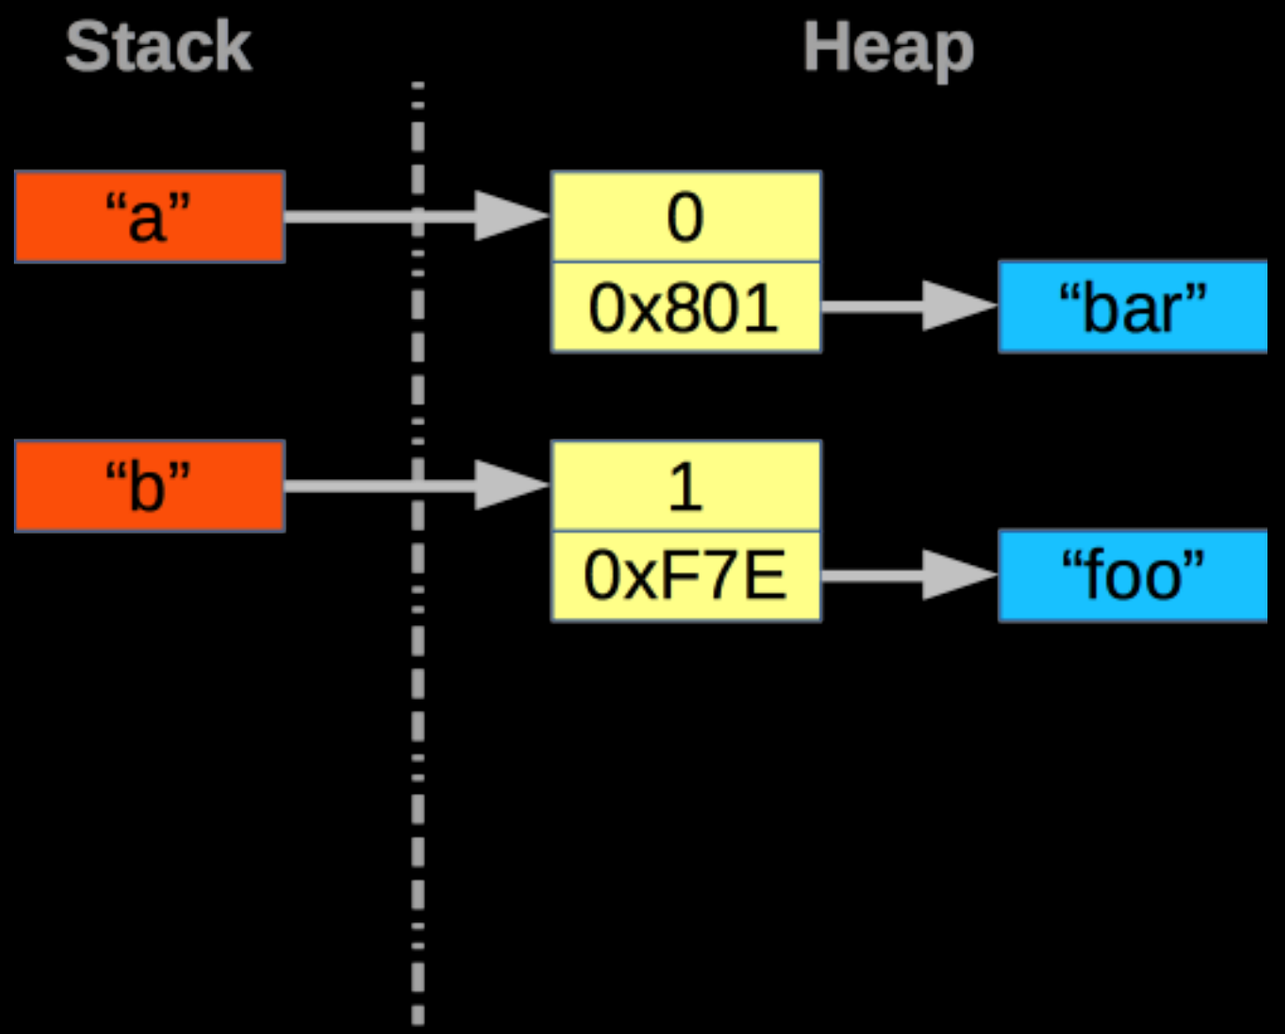

Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
>>> del a
```

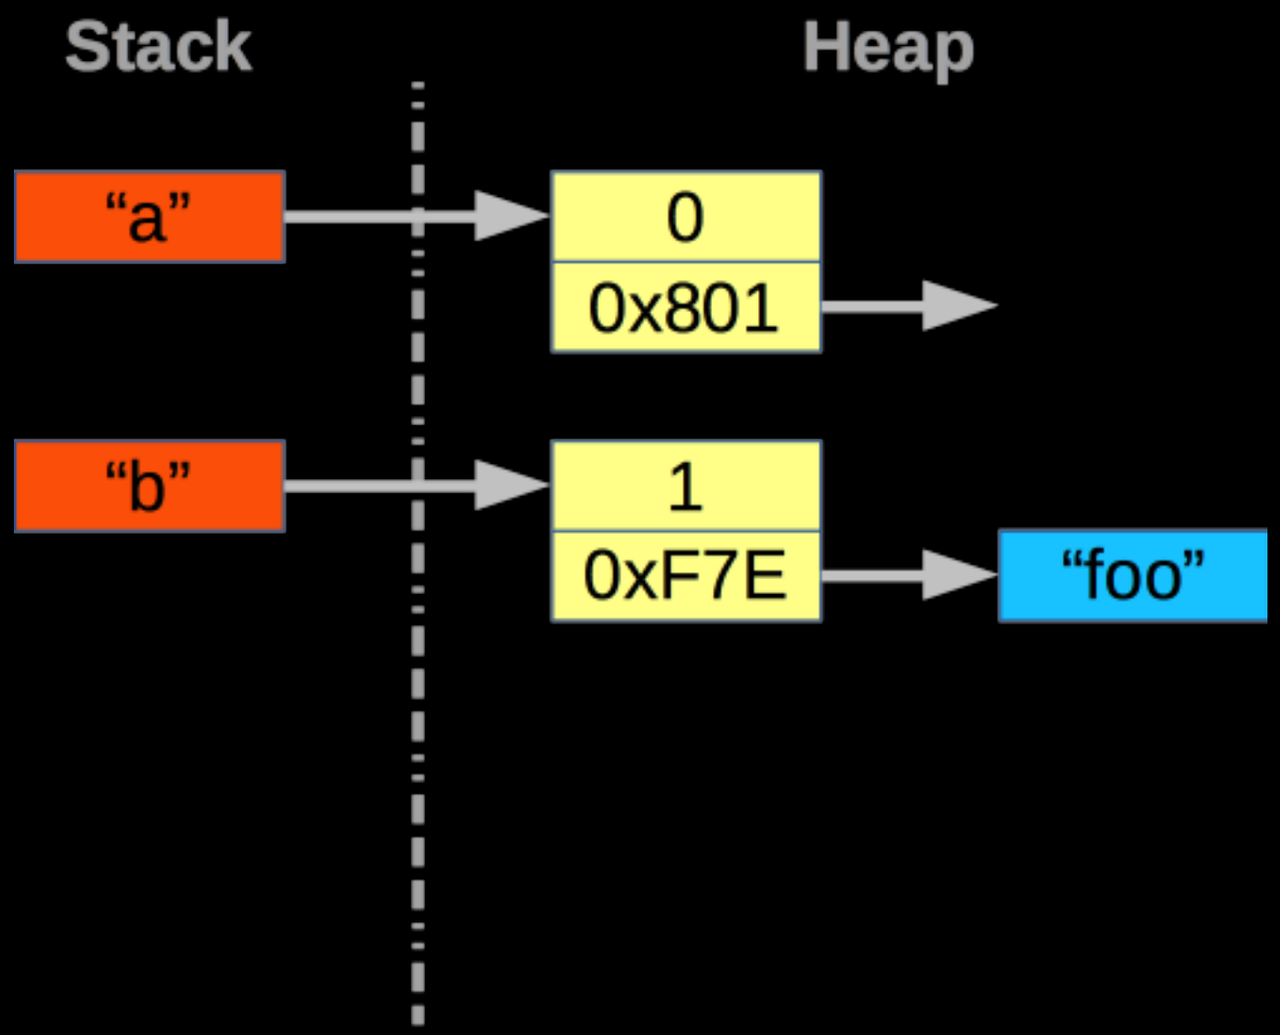

Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
>>> del a
```

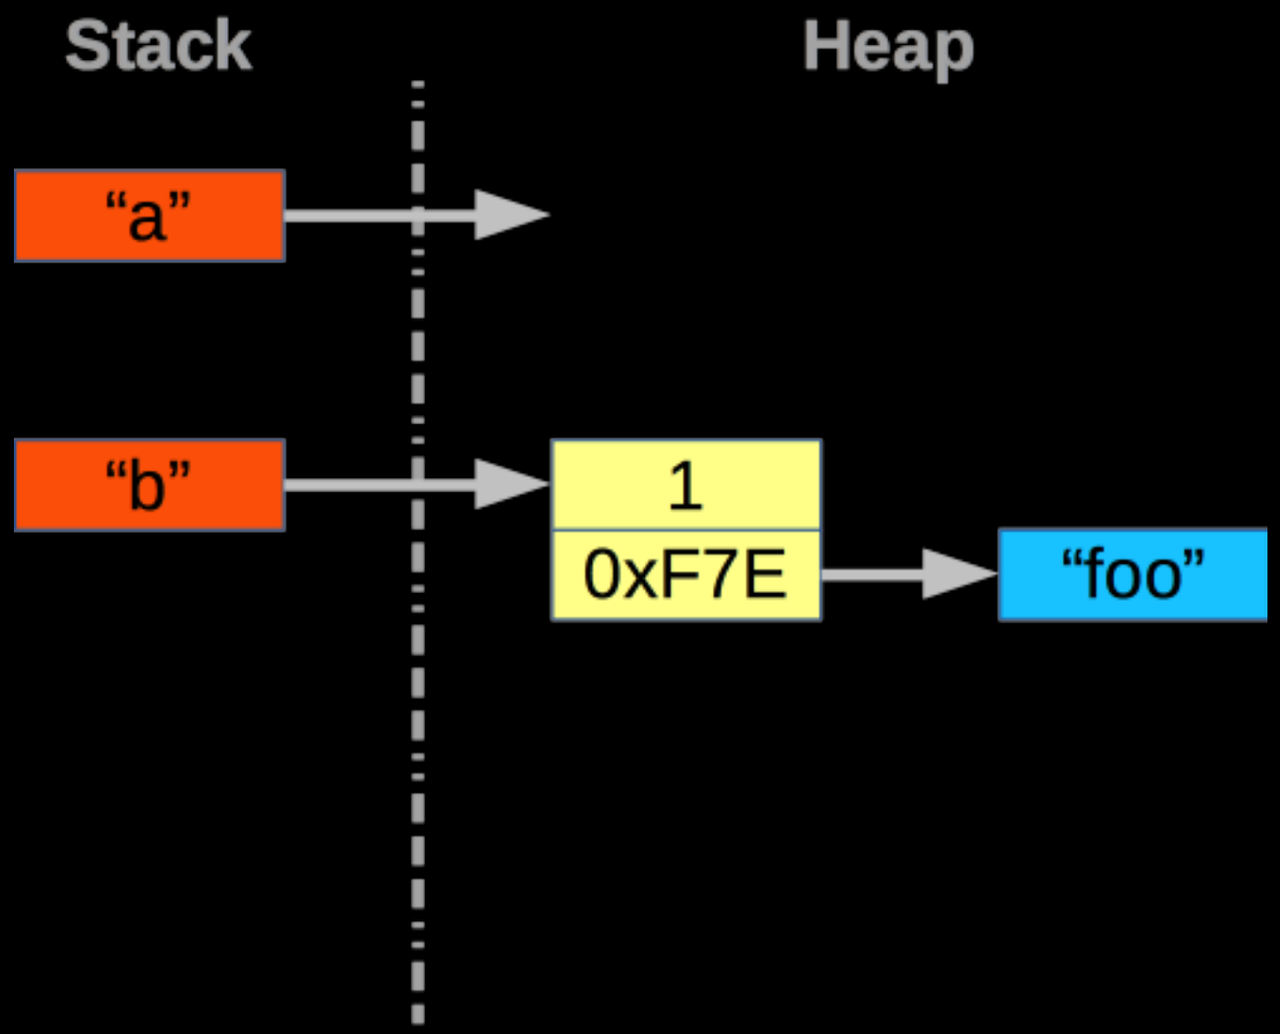

Reference Counting

```
>>> a = 'foo'
>>> b = a
>>> a = 'bar'
>>> del a
```

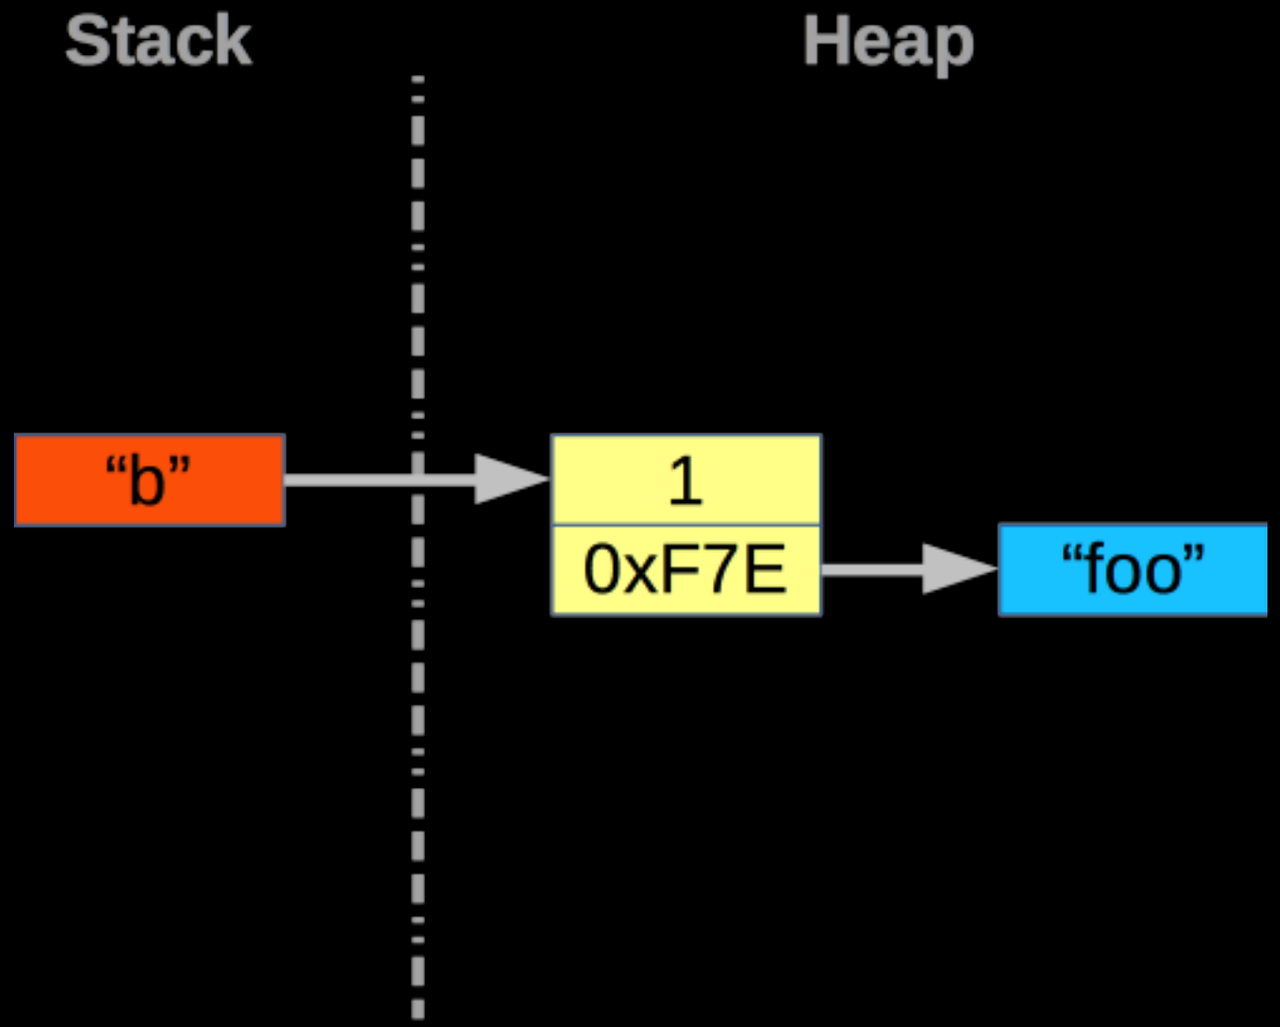

Garbage Collection

- The GC is just there to resolve cyclic references
  - Only works with containers
- It is not a Unicorn
  - Will not reclaim lost C allocated memory
  - Will not reclaim lost PyObject references

Summary

1 coding pattern to keep the dragons at bay

2 things to avoid

3 kinds of `PyObject*` used in CPython

Summary

1 coding pattern to keep the dragons at bay

2 things to avoid

3 kinds of `PyObject*` used in CPython

2 Things to Avoid

- Memory leaks
- Access after deallocation

C Memory Leaks

```c
void leak() {
    char *p = malloc(1024);
    p[8] = 'A';
}
```

C Access After `free()`

["free" is **not** "make impossible to access"]

```c
void access_after_free() {
    char *p = malloc(1024);
    p[8] = 'A';

    free(p);
    printf("%c", p[8]);
}
```

Py Memory Leaks

```c
#include "Python.h"

void py_leak() {
    PyObject *pObj;
    pObj = PyBytes_FromString("Hello world\n");
    PyObject_Print(pObj, stdout, 0);


}
```

Py Memory Leaks - Fixed

```
#include "Python.h"

void py_leak() {
    PyObject *pObj;
    pObj = PyBytes_FromString("Hello world\n");
    PyObject_Print(pObj, stdout, 0);
    Py_DECREF(pObj);

}
```

Py Access After DecRef

```
#include "Python.h"

void py_access_after_free() {
    PyObject *pObj;
    pObj = PyBytes_FromString("Hello world\n");
    PyObject_Print(pObj, stdout, 0);
    Py_DECREF(pObj);
    PyObject_Print(pObj, stdout, 0);
}
```

Py Access After DecRef

Please don't do this

```
Py_DECREF(pObj);

/* Is ob_refcnt really for the same object? */
if (pObj->ob_refcnt > 0) {
    PyObject_Print(pObj, stdout, 0);
}
```

Summary

1 coding pattern to keep the dragons at bay

2 things to avoid

3 kinds of `PyObject*` used in CPython

3 Reference Types

- **New** references occur when a `PyObject` is created
  - Example: creating a new list.
- **Stolen** references occur when a `PyObject` is created and assigned. Typically 'setters'
  - Example: appending a new value to a list.
- **Borrowed** references are used when getting a `PyObject`
  - Example: accessing a member of a list.
  - If ***shared*** references mean more to you, great! Thats exactly what they are.

Programming by Contract

- **New** `PyObject*` Your job to deallocate it
  - Or give it to someone who will
- **Stolen** `PyObject*` The 'thief' will deallocate it
  - Do not do so yourself
- **Borrowed** `PyObject*` The real owner can deallocate it at any time
  - Unless you prevent them by registering your interest

New References

```
static PyObject *subtract_long(long a, long b) {
    PyObject *pA, *pB, *r;

    pA = PyLong_FromLong(a);          /* New ref */
    pB = PyLong_FromLong(b);          /* New ref */

    r = PyNumber_Subtract(pA, pB);  /* New ref */

    Py_DECREF(pA);            /* I must decref */
    Py_DECREF(pB);            /* I must decref */
    return r;                 /* Caller must decref */
}
```

New References
Please Don't do this

```
static PyObject *subtract_long(long a, long b) {
    return PyNumber_Subtract(
                /* A leak */
                PyLong_FromLong(a),
                /* Another leak */
                PyLong_FromLong(b)
                );
}
```

Stolen References

```
static PyObject *make_tuple() {
    PyObject *r, *v;

    r = PyTuple_New(3);          /* New ref */
    v = PyLong_FromLong(1L);     /* New ref */
    PyTuple_SetItem(r, 0, v);

    v = PyLong_FromLong(2L);     /* New ref */
    PyTuple_SetItem(r, 1, v);

    /* More common pattern */
    PyTuple_SetItem(r, 2, PyLong_FromLong(3L));
    return r;    /* Callers must decref */
}
```

Stolen References

Please don't do this

```
PyObject *r, *v;

r = PyTuple_New(3);             /* New ref */

v = PyLong_FromLong(1L);        /* New ref */
PyTuple_SetItem(r, 0, v);       /* r 'steals' v */

Py_DECREF(v); /* NO! v 'belongs' to r */
```

'Borrowed' References

- These are generally 'getters'

```
PyObject *pList = ...
PyObject *pVal  = PyList_GetItem(pList, 0);
```

'Borrowed' References

- These are generally 'getters'

```
PyObject *pList = ...
PyObject *pVal  = PyList_GetItem(pList, 0);
```

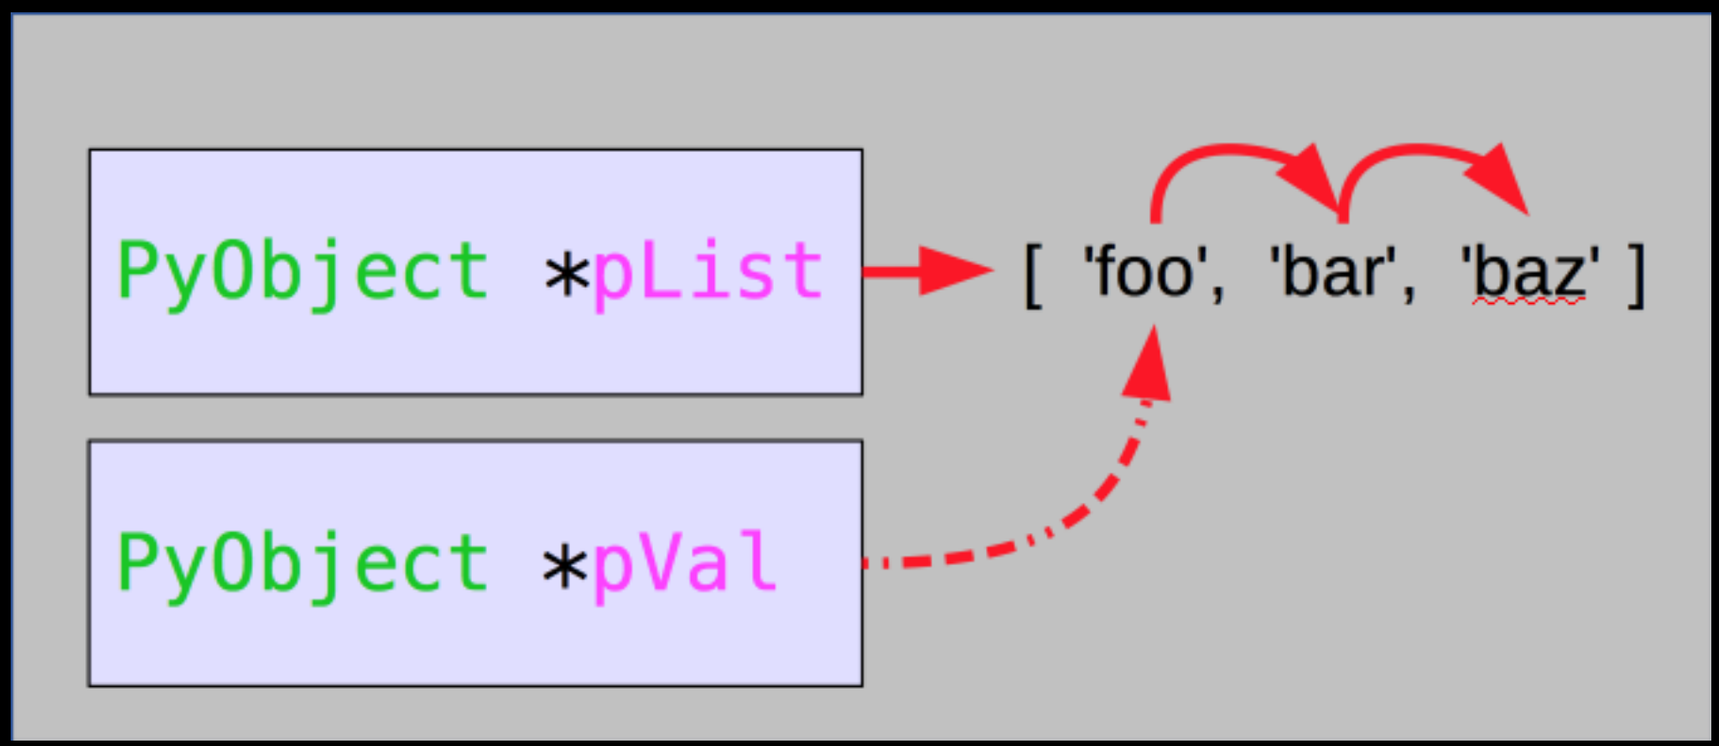

'Borrowed' References

- Multiple pointers to the same object - Aaargh!
  - Which is responsible for deallocating the object?
  - What happens to the other pointers when one deallocates the object?
- They can be the source of the most subtle bugs

'Borrowed' References

```
static PyObject *
borrow_BAD(PyObject *pList) {
    PyObject *pFirst;
    pFirst = PyList_GetItem(pList, 0);

    function(pList);    /* Dragons ahoy! */
    PyObject_Print(pFirst, stdout, 0);


    Py_RETURN_NONE;
}
```

Hmm… Suppose

```
static PyObject *
borrow_BAD(PyObject *pList) {
    PyObject *pFirst;
    pFirst = PyList_GetItem(pList, 0);

    function(pList);    /* Dragons ahoy! */
    PyObject_Print(pFirst, stdout, 0);


    Py_RETURN_NONE;
}
```

Hmm… Suppose

```
static PyObject *
borrow_BAD(PyObject *pList) {
    PyObject *pFirst;
    pFirst = PyList_GetItem(pList, 0);

    function(pList);    /* Dragons ahoy! */
    PyObject_Print(pFirst, stdout, 0);

    Py_RETURN_NONE;
}
```

This removed the first item in the list!

Borrowed Ref Dragon 0

```
>>> import cPyRefs
>>> l = ['foo', 'bar', 'baz']
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

Borrowed Ref Dragon 1

```
>>> import cPyRefs
>>> l = ['foo', 'bar', 'baz']
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

```
>>> import cPyRefs
>>> l = ['foo', 'bar', 'baz']
>>> a = l[0]
>>> cPyRefs.borrow_bad(l) # Works fine!
```

Borrowed Ref Dragon 1

```
>>> import cPyRefs
>>> l = ['foo', 'bar', 'baz']
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

```
>>> import cPyRefs
>>> x ='foo'
>>> l = ['bar', 'baz']
>>> l.insert(0, x)
>>> cPyRefs.borrow_bad(l) # Works fine!
```

Borrowed Ref Dragon 2

```
>>> import cPyRefs
>>> l = ['foo', 'bar', 'baz']
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

Borrowed Ref Dragon 2

```
>>> import cPyRefs
>>> l = ['foo', 'bar', 'baz']
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

```
>>> import cPyRefs
>>> l = [1, 2, 3]
>>> cPyRefs.borrow_bad(l) # Works fine!
```

Borrowed Ref Dragon 2

```
>>> import cPyRefs
>>> l = ['foo', 'bar', 'baz']
>>> cPyRefs.borrow_bad(l) # SEGFAULT!
```

```
>>> import cPyRefs
>>> l = [1, 2, 3]
>>> cPyRefs.borrow_bad(l) # Works fine

>>> import cPyRefs
>>> l = [800, 801, 802]
>>> cPyRefs.borrow_bad(l) # Kaboom!
```

Run-time Errors
+ Data Dependent Errors

Run-time Errors
+ Data Dependent Errors

The Problem

```
static PyObject *
borrow_BAD(PyObject *pList) {
    PyObject *pFirst;
    pFirst = PyList_GetItem(pList, 0);

    function(pList);    /* Dragons ahoy! */
    PyObject_Print(pFirst, stdout, 0);


    Py_RETURN_NONE;
}
```

The Fix

```
static PyObject *
borrow_BAD(PyObject *pList) {
    PyObject *pFirst;
    pFirst = PyList_GetItem(pList, 0);

    function(pList);    /* Dragons ahoy! */
    PyObject_Print(pFirst, stdout, 0);


    Py_RETURN_NONE;
}
```

The Fix

```
static PyObject *
borrow_BAD(PyObject *pList) {
    PyObject *pFirst;
    pFirst = PyList_GetItem(pL
    Py_INCREF(pFirst);
    function(pList);     /* Dragons lurked! */
    PyObject_Print(pLast, stdout, 0);
    Py_DECREF(pFirst);
    pFirst = NULL;
    Py_RETURN_NONE;
}
```

Register your interest!

Let go

Summary

1 coding pattern to keep the dragons at bay

2 things to avoid

3 kinds of `PyObject*` used in CPython

1 Pattern For Reliable C

- Borrowed references incref'd and decref'd correctly.
- A single place for clean up code
  - No early returns
- Exception consistency. Either:
  - An exception is set **and** NULL is returned.
  - Or: no Exception set **and** non-NULL returned.

Writing Pythonic Python

```
def function(obj):
    ret = None;
    try:
        # Do fabulous stuff here
        # On error, raise
    except ... as err:
        # Handle exceptions
    finally:
        # And we are out
    return ret;
```

Writing Pythonic C

Writing Pythonic C

```c
static PyObject *function(PyObject *arg1) {
    PyObject *ret = NULL;

    goto try;
try:
    /* Do fabulous stuff here */
    /* On error "goto except;" */
    goto finally;
except:
    /* Handle exceptions */
finally:
    /* And we are out */
    return ret;
}
```

Function Entry

```
static PyObject *function(PyObject *arg1) {
    /* Create any local PyObject* as NULL */
    PyObject *obj_a    = NULL;
    /* Create the PyObject* return value as NULL */
    PyObject *ret      = NULL;

    goto try; /* Pythonic 'C' ;-) */
try:
```

try:

```
try:
    assert(! PyErr_Occurred());
    /* Inc the reference count of the arguments. */
    assert(arg1);
    Py_INCREF(arg1);

    /* Your code here */

    /* Local object creation; borrowed or new. */
    obj_a = ...;
    /* If an error ... */
    if (! obj_a) {
        PyErr_SetString(PyExc_ValueError, "Ooops.");
        goto except;
    }
```

try:

```
/* Return object creation, ret will either be a
 * new reference or a borrowed reference
 * INCREF'd */
ret = ...;
if (! ret) {
        PyErr_SetString(PyExc_ValueError,
                        "Ooops again.");
        goto except;
    }
    /* If success then check exception is clear,
     * goto finally; with non-NULL return value. */
    assert(! PyErr_Occurred());
    assert(ret);
    goto finally;
except:
```

except:

```
except:
    /* Failure so Py_XDECREF the return value */
    Py_XDECREF(ret);
    /* Check a Python error set somewhere above */
    assert(PyErr_Occurred());
    /* Signal failure */
    ret = NULL;
    /* Fall through to finally: */
finally:
```

finally:

```
finally:
    /* All _local_ PyObjects are Py_XDECREF'd here.
     * For new references this will free them.
     * For borrowed references this
     * will return them to their previous state. */
    Py_XDECREF(obj_a);
    /* Decrement the ref count of given arguments
     * if they have been incremented. */
    Py_DECREF(arg1);
    /* And return... */
    return ret;
}
```

All this and more…

https://github.com/paulross

In “PythonExtensionPatterns”

The Documentation is Excellent - Use it!

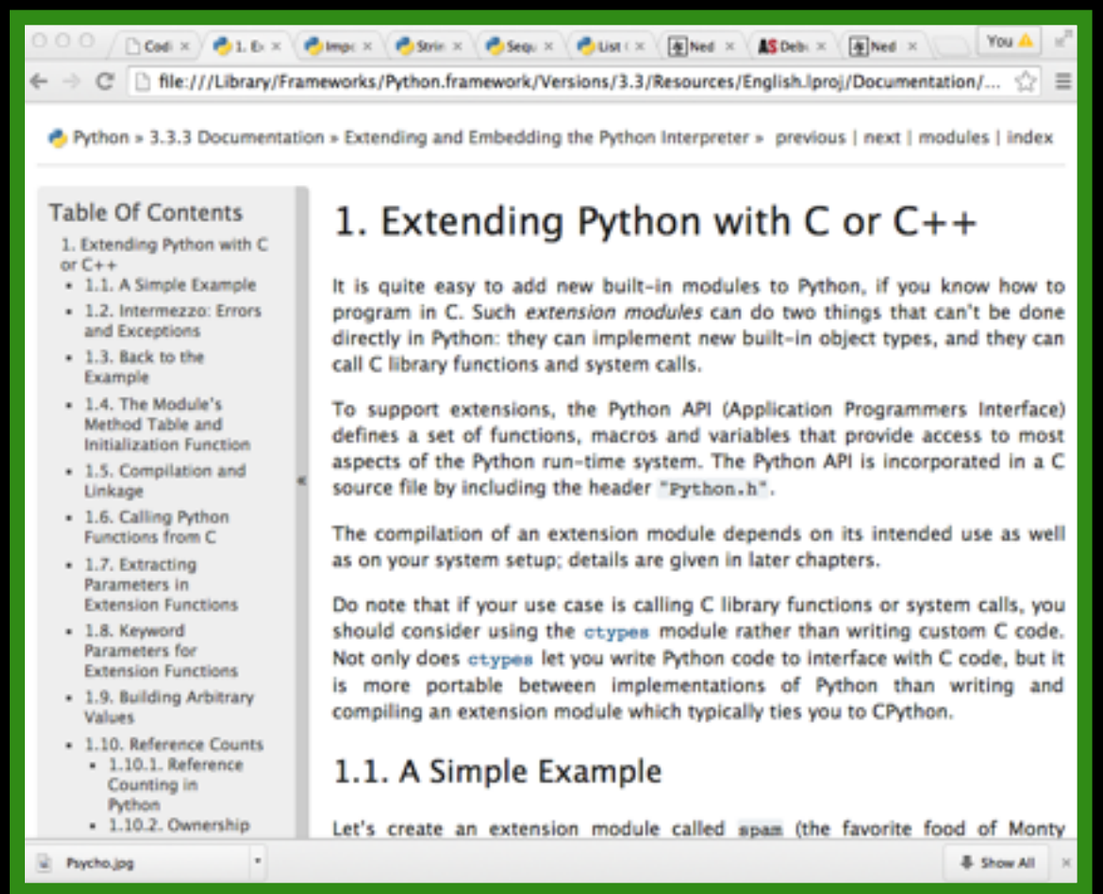

The Documentation is Excellent - Use it!

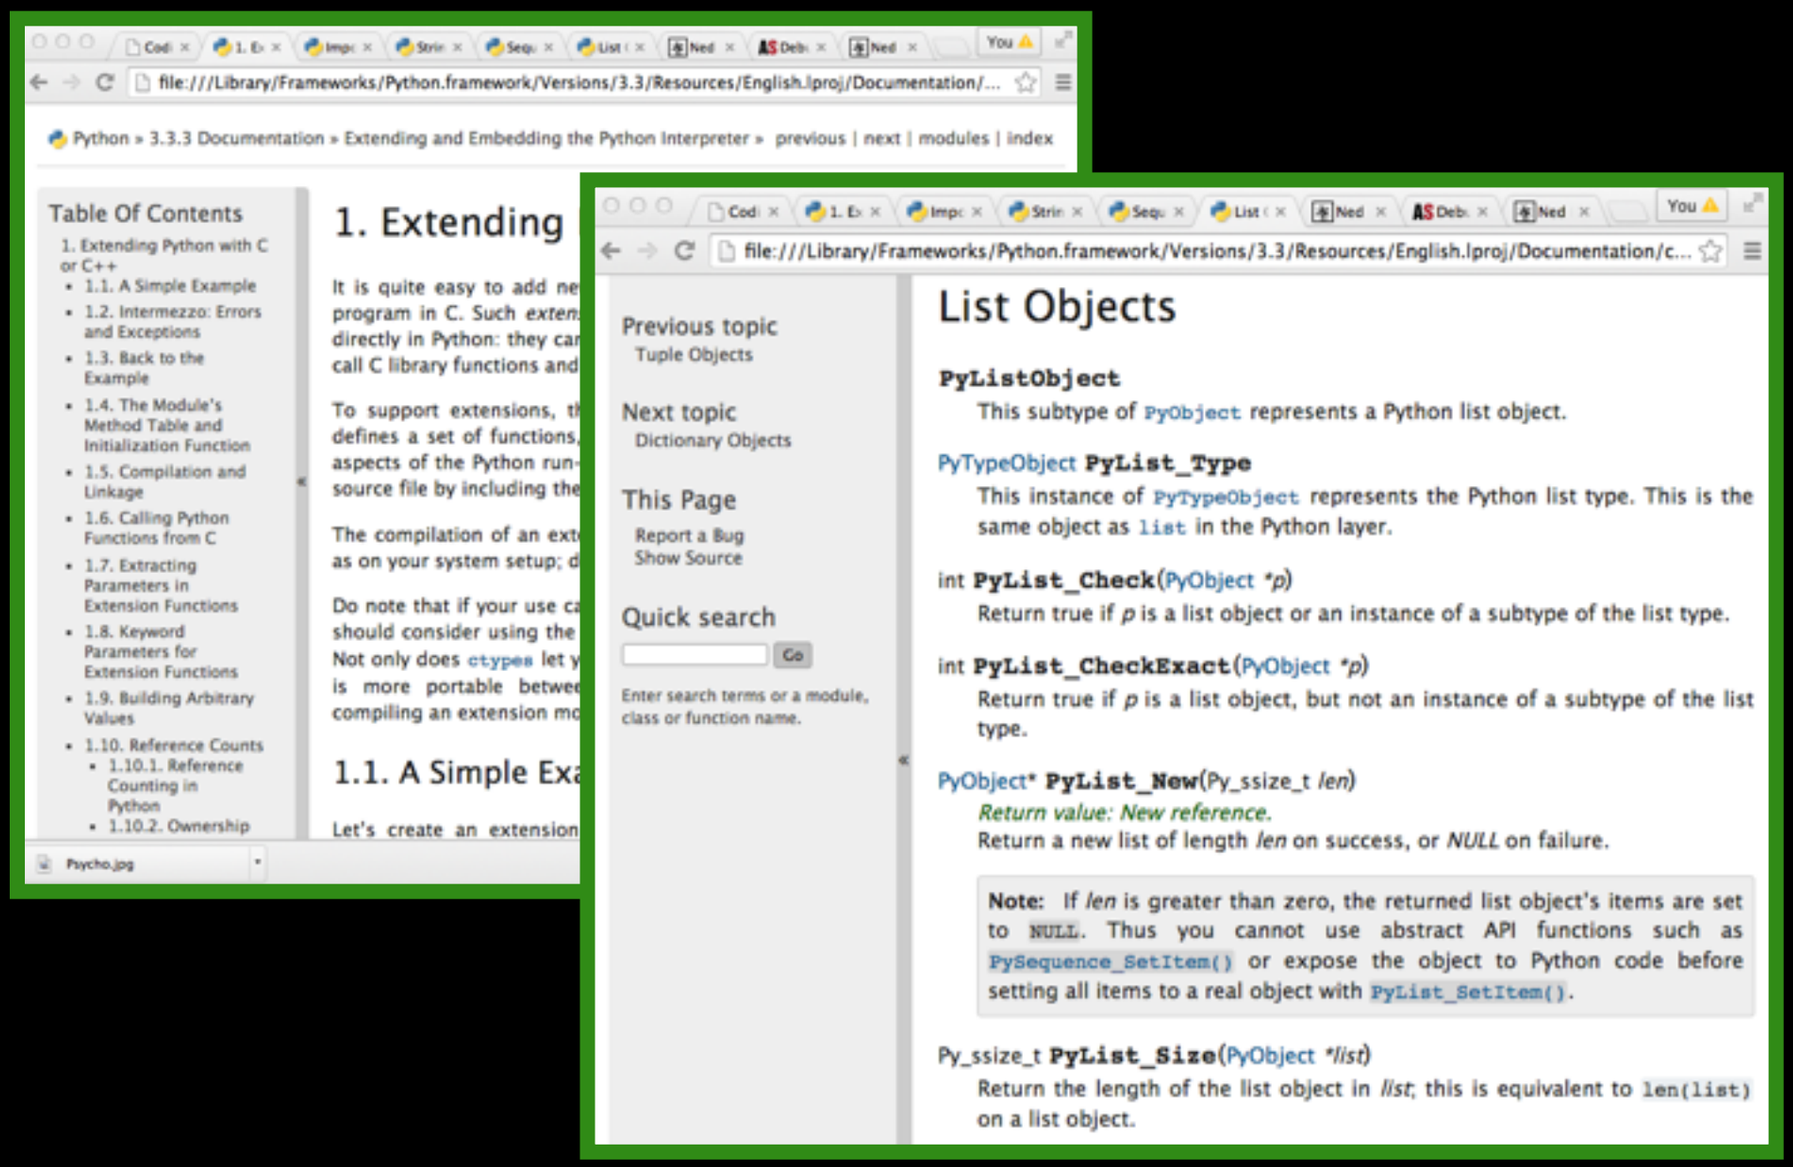

The Documentation is Excellent - Use it!

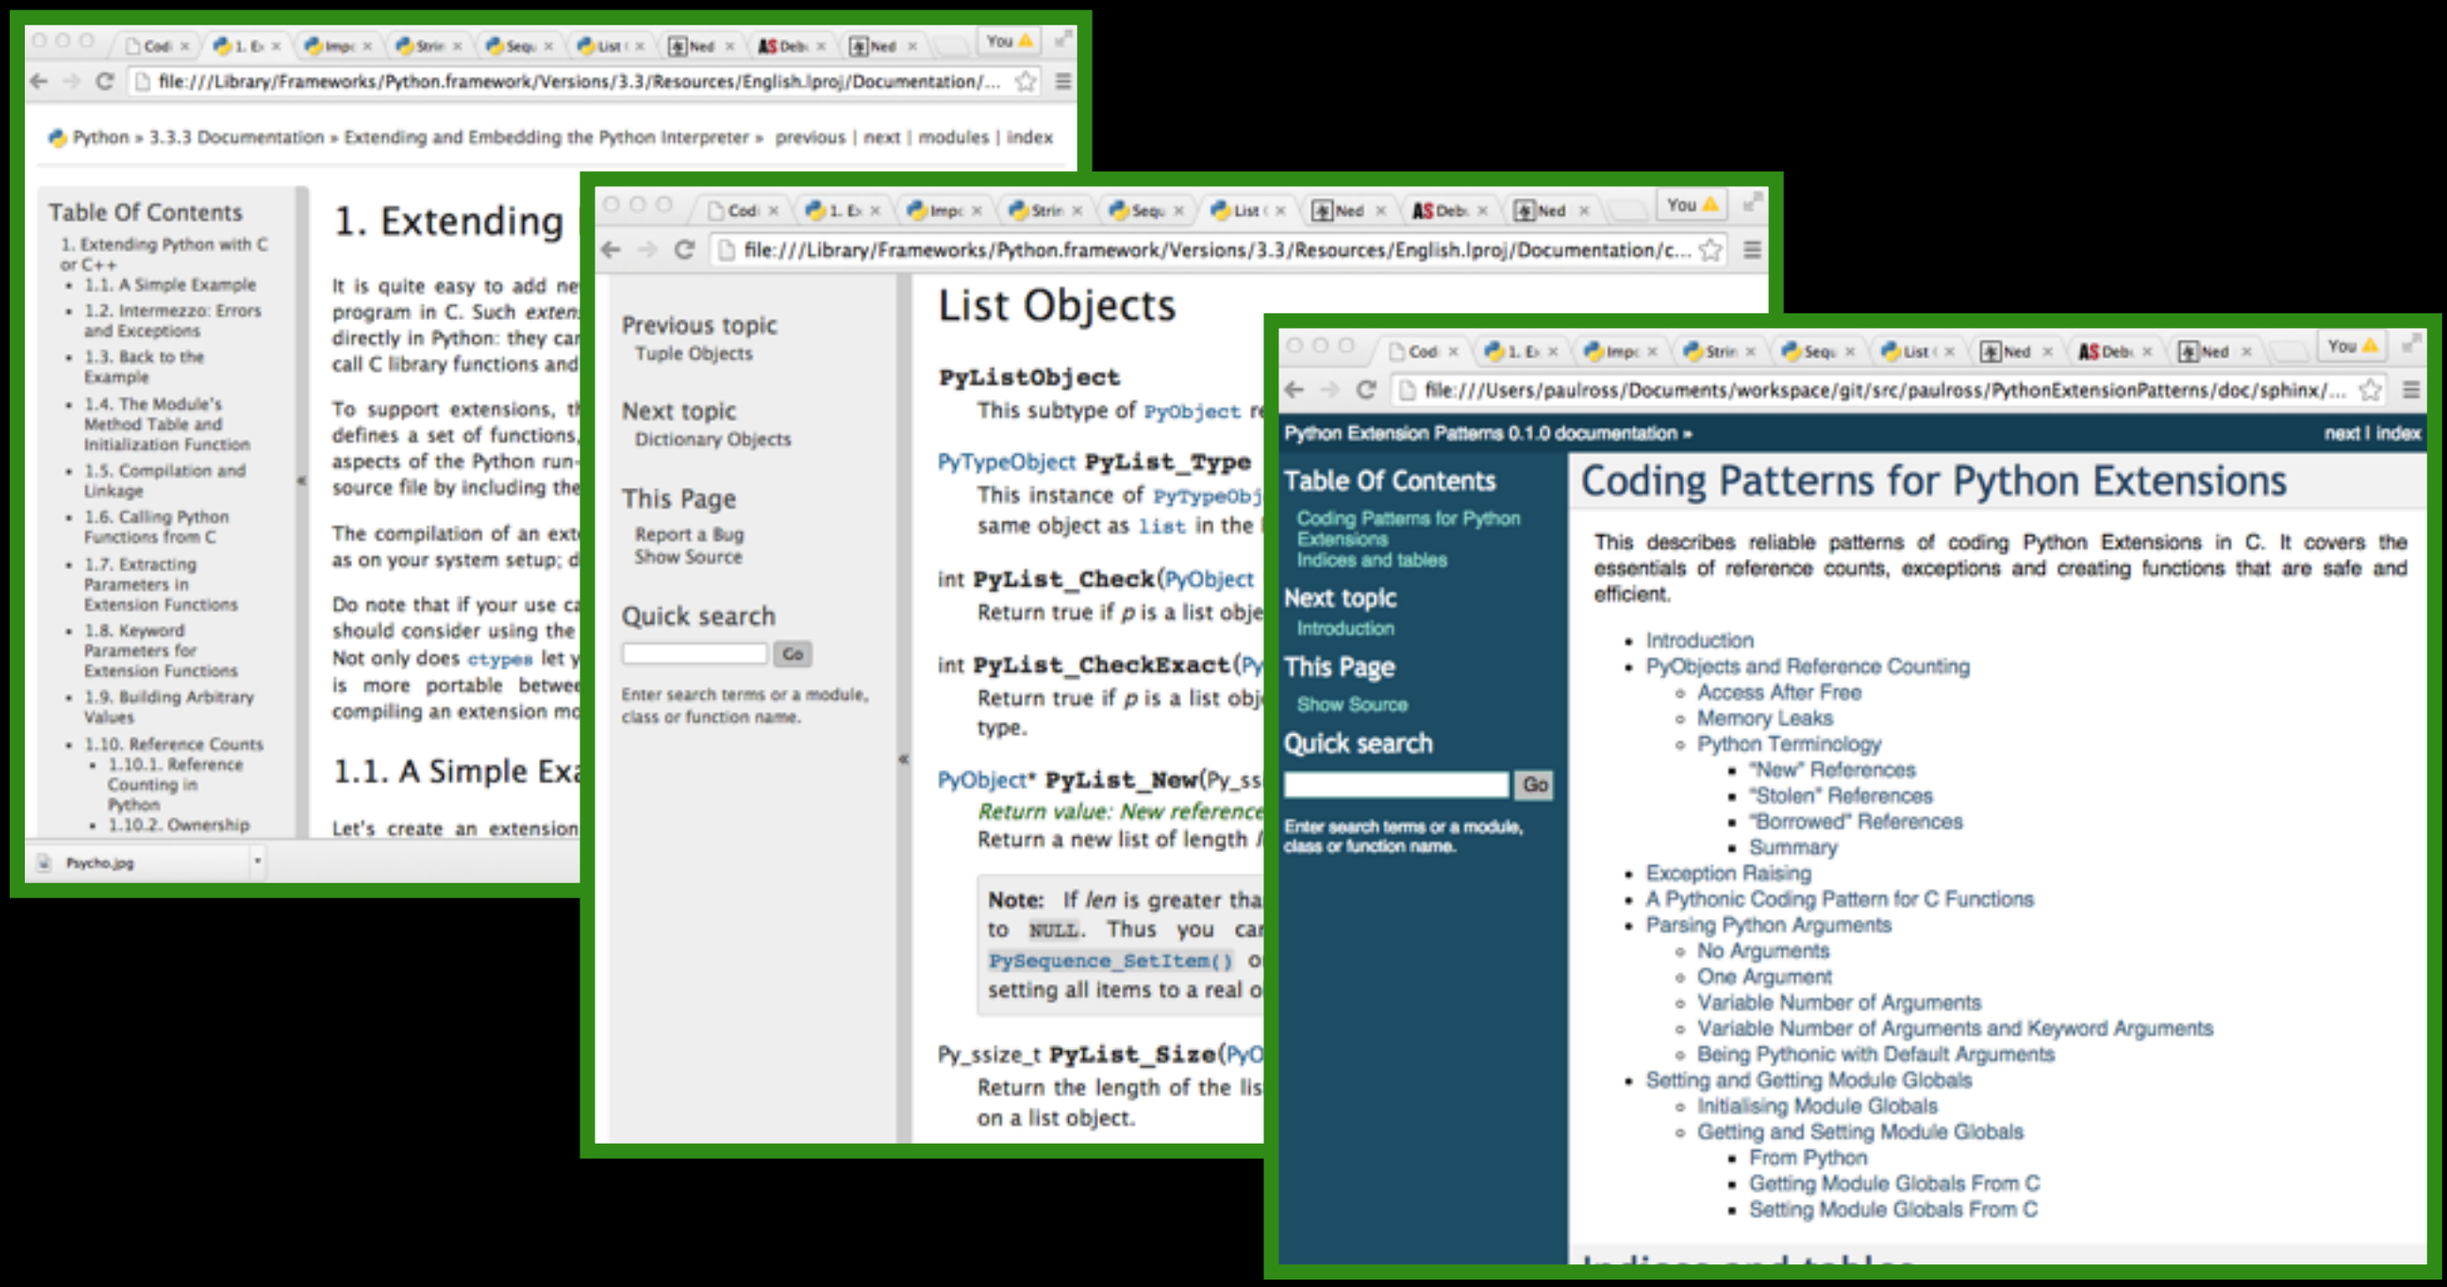

https://github.com/paulross

War Story ~ Mandatory

- Proprietary binary files of oilfield data
  - Self describing, variable format, sequentially written
- Make them random access by creating an index
  - The index is built with a sequence of `seek()`/`read()` operations
  - `read()` is about 1% to 2% of the original file size
- Originally written in Python. Typ. 10-100ms/Mb
- How fast can we go?

Gnuplot

Improvement in indexing cost using C extension cPhysRec.

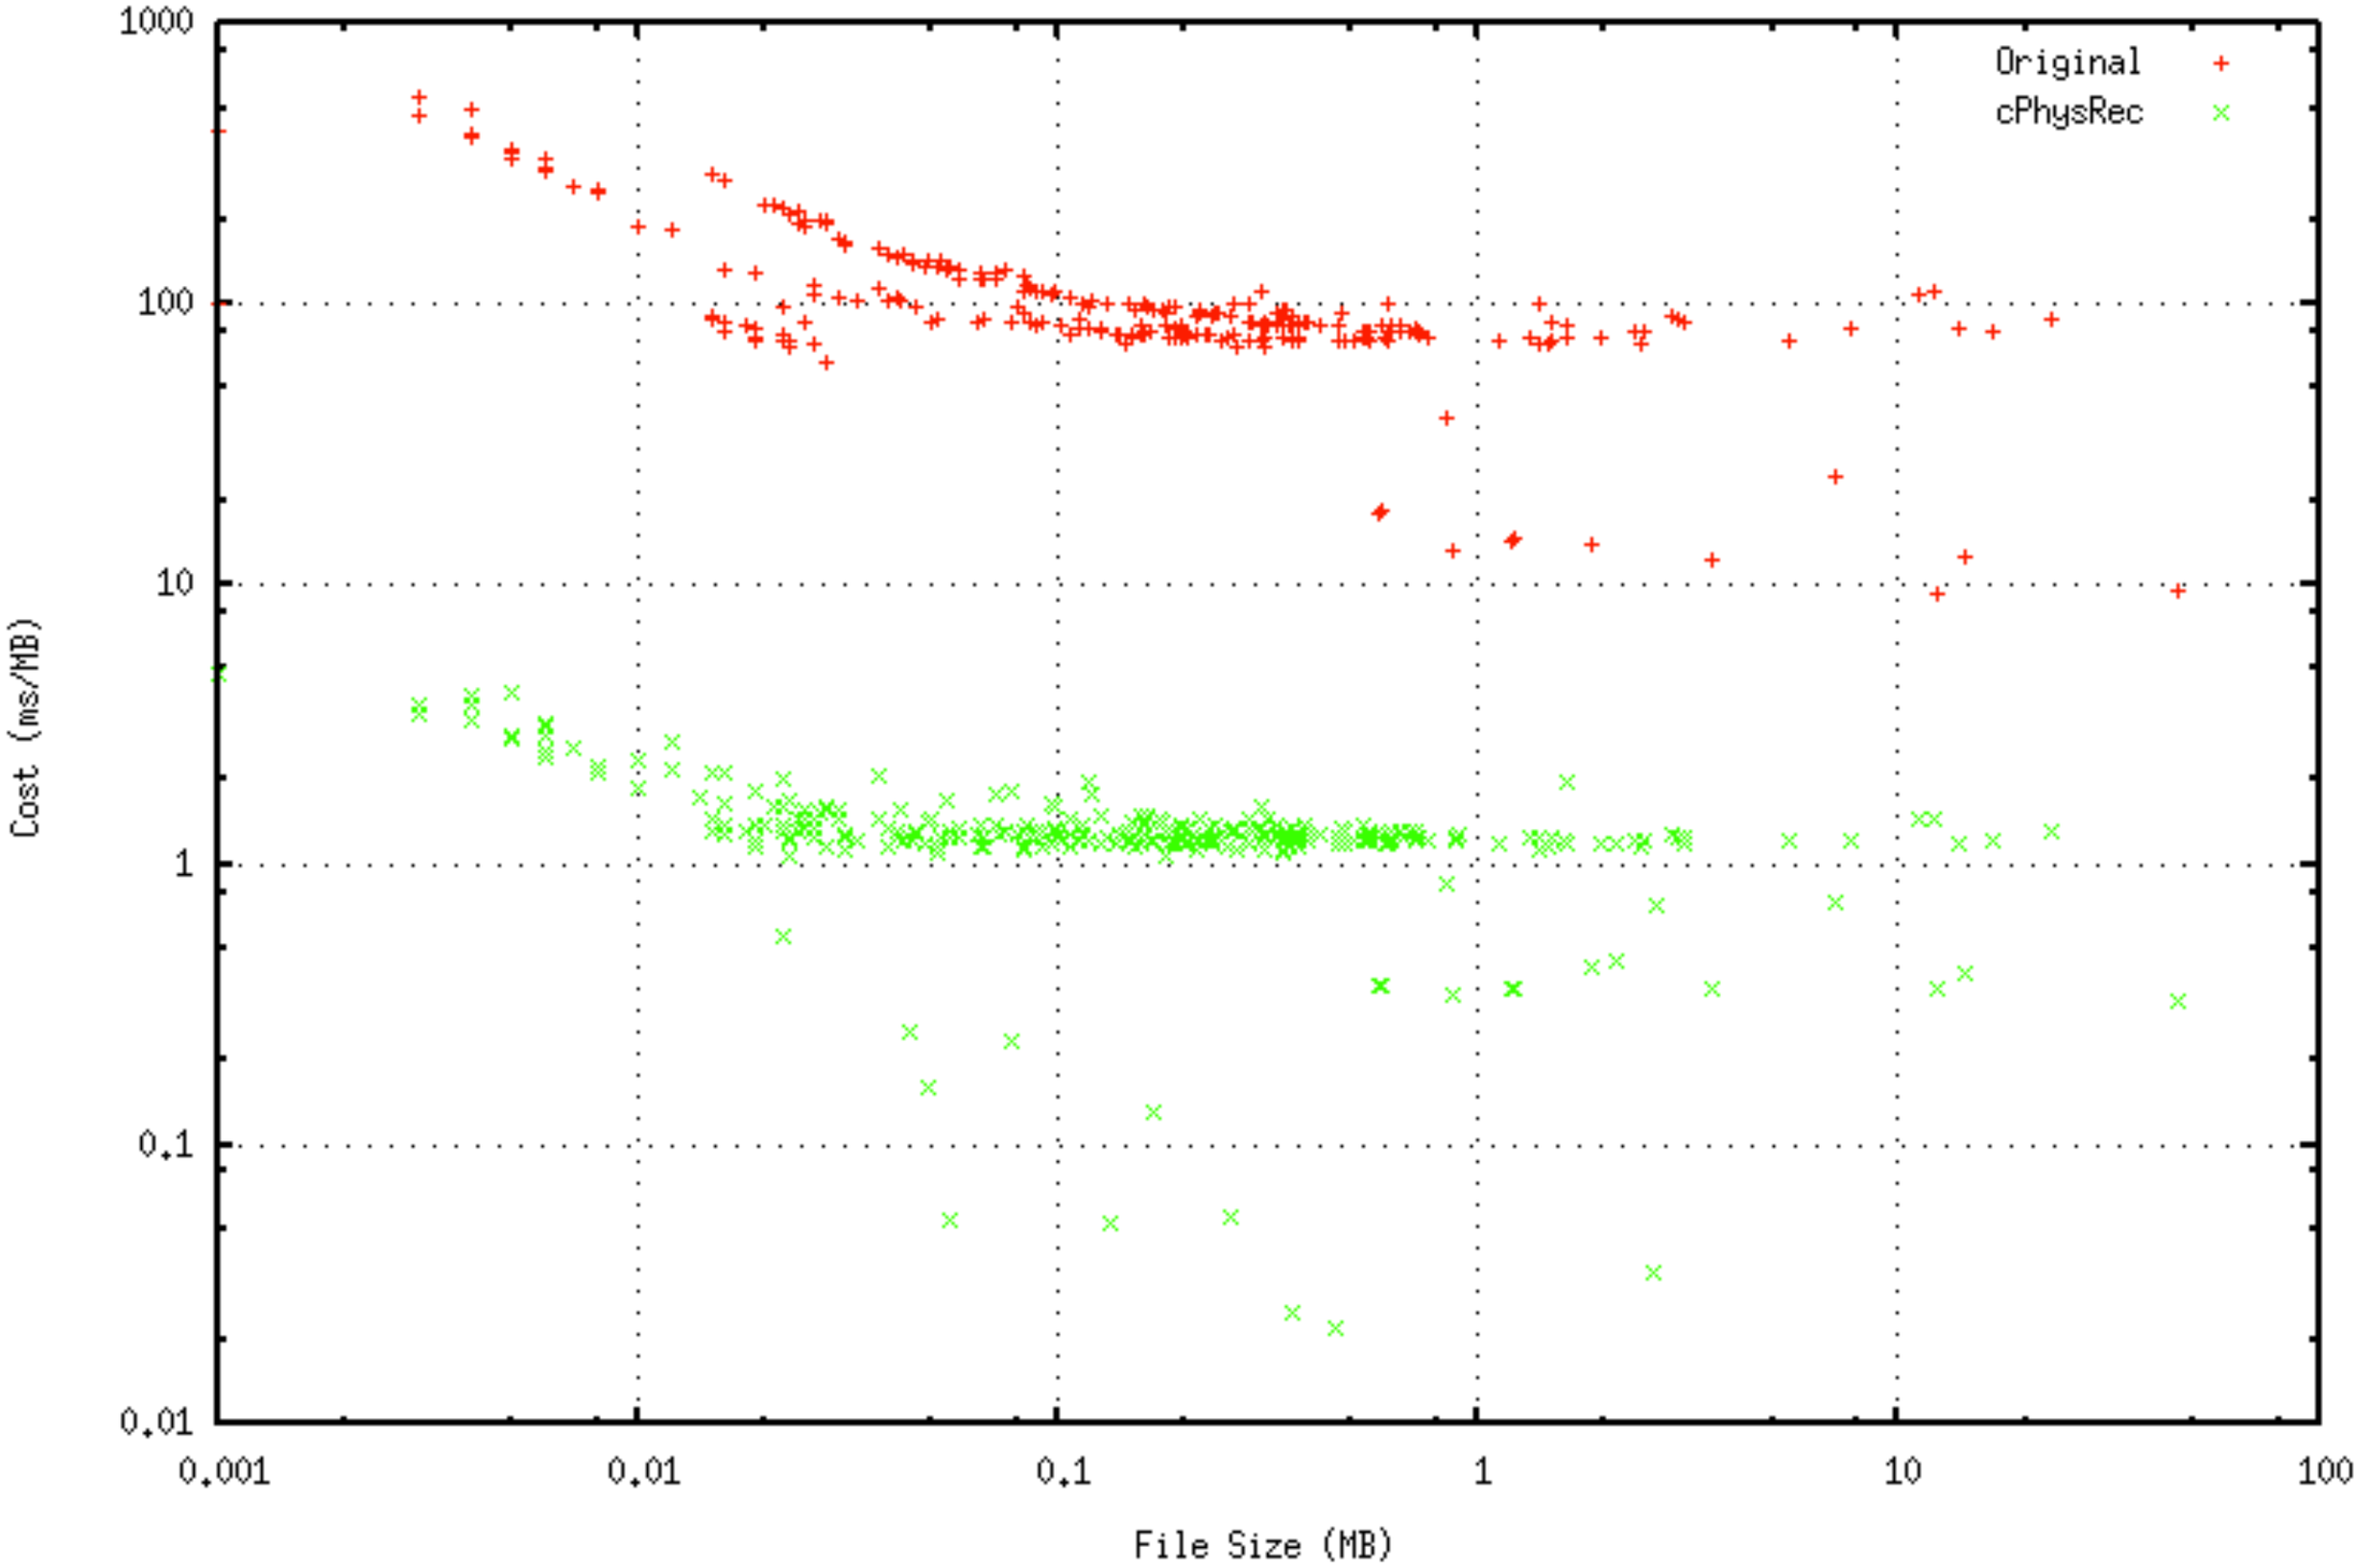

0.000272545, 0.00274264

Summary

1 coding pattern to keep the dragons at bay

2 things to avoid

- Allocation with no deallocation
- Access after deallocation

3 kinds of references to `PyObject*`

- **New**: its yours
- **Stolen**: its theirs
- **Borrowed**: your sharing something that is theirs - let them know!

github.com/paulross

twitter.com/manahltech

Questions ???

github.com/paulross

twitter.com/manahltech